

Convolutional Neural Network

Image Classification using CNN

Group Members:

Mehar Abdul Wahab (201980053)

Mehar Hamid Ishfaq (201980038)

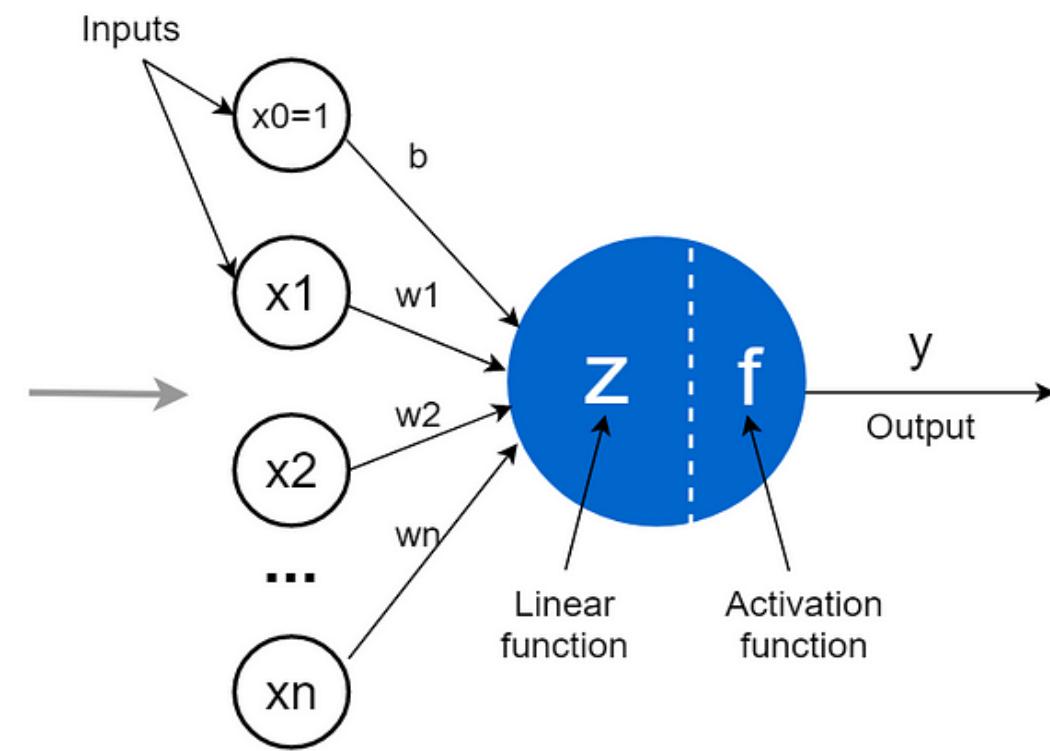
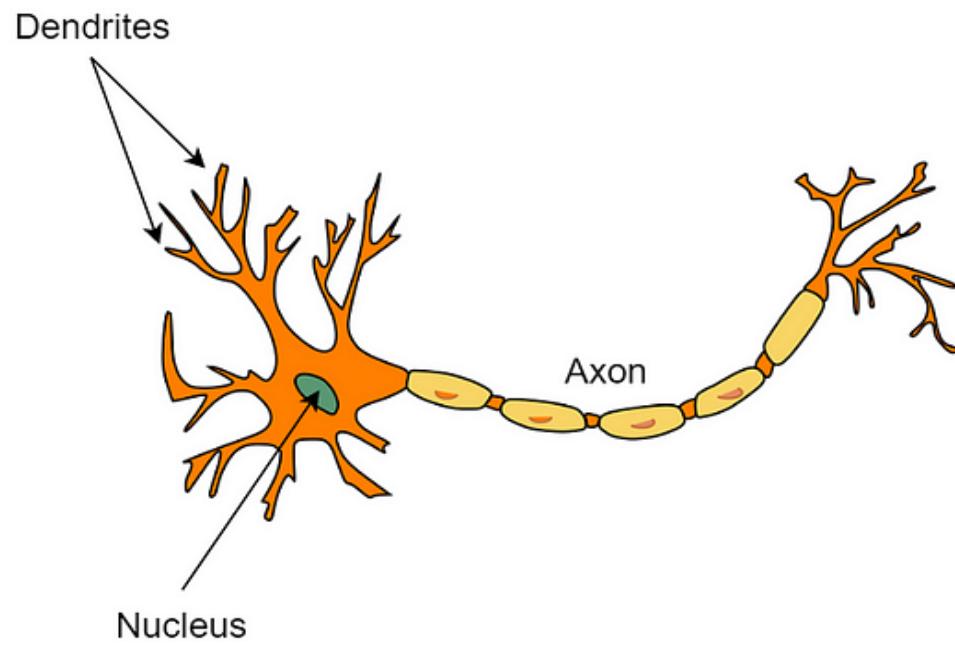
Rajpoot Sufyan Ahmad (201980013)

Mughal Nadeem (201980050)

Applications

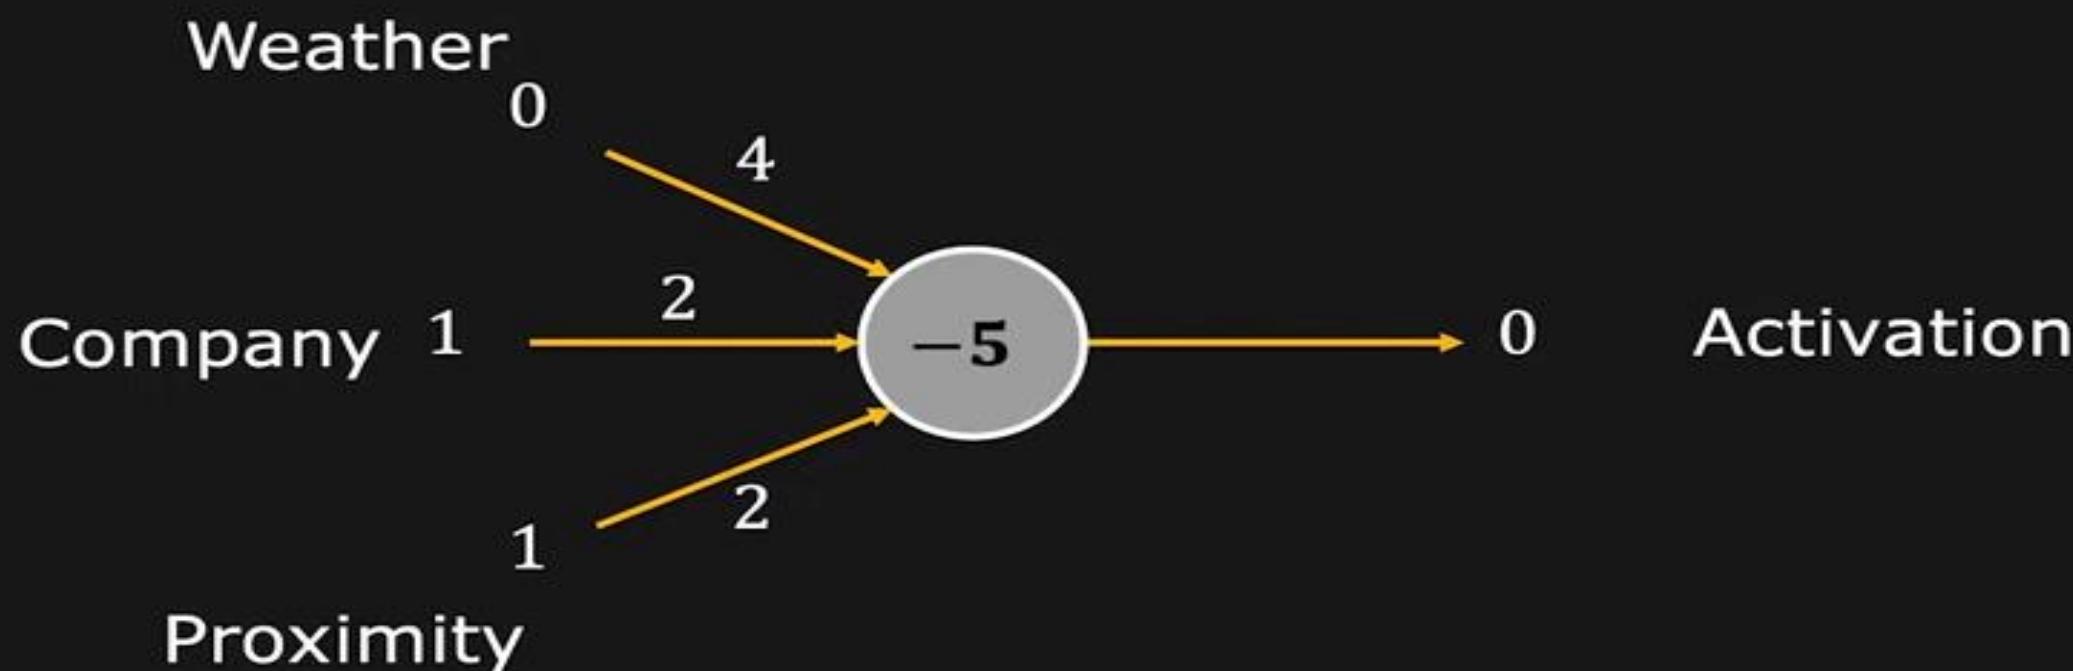
- Object Recognition and Classification
- Quality Control in Manufacturing
- Agriculture
- Food Recognition
- Social Media
- Gesture Recognition
- Self-Driving Cars

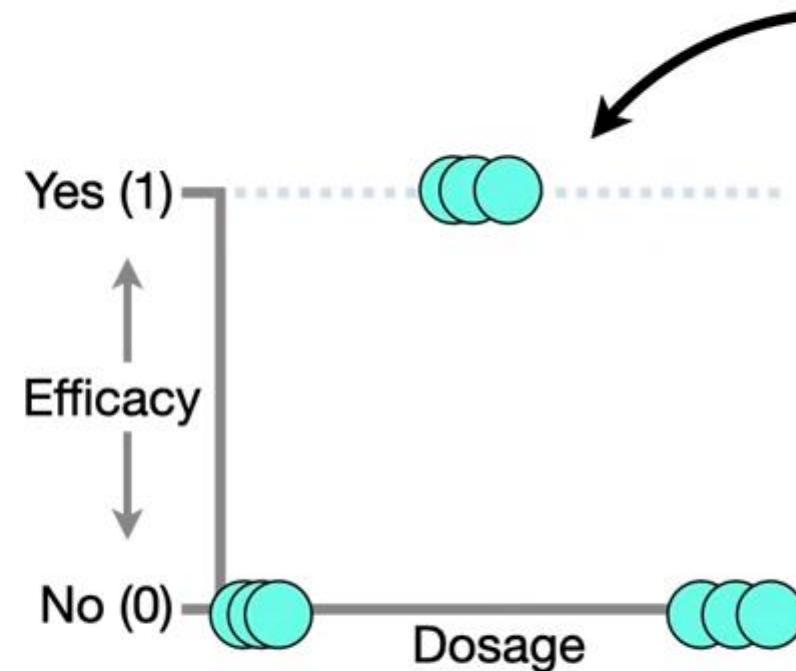
Biological vs Artificial Neuron



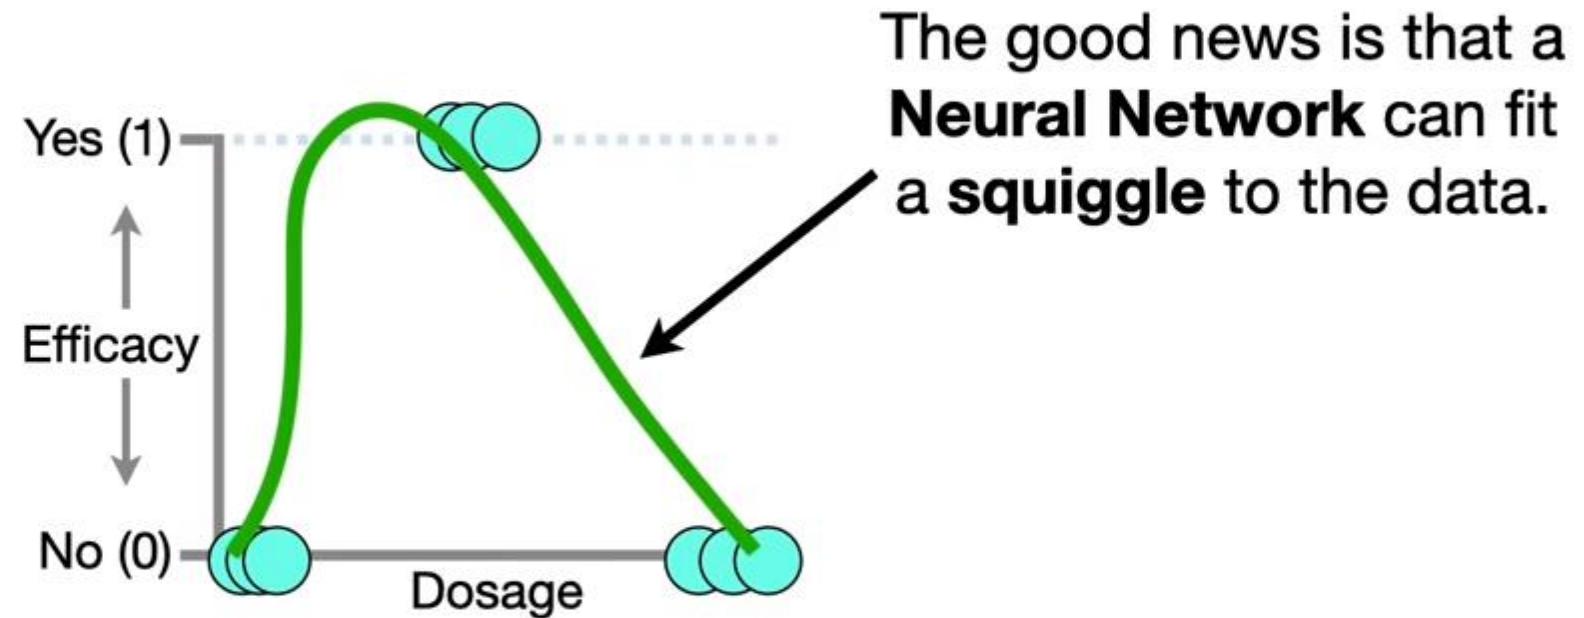
Perceptron: Example

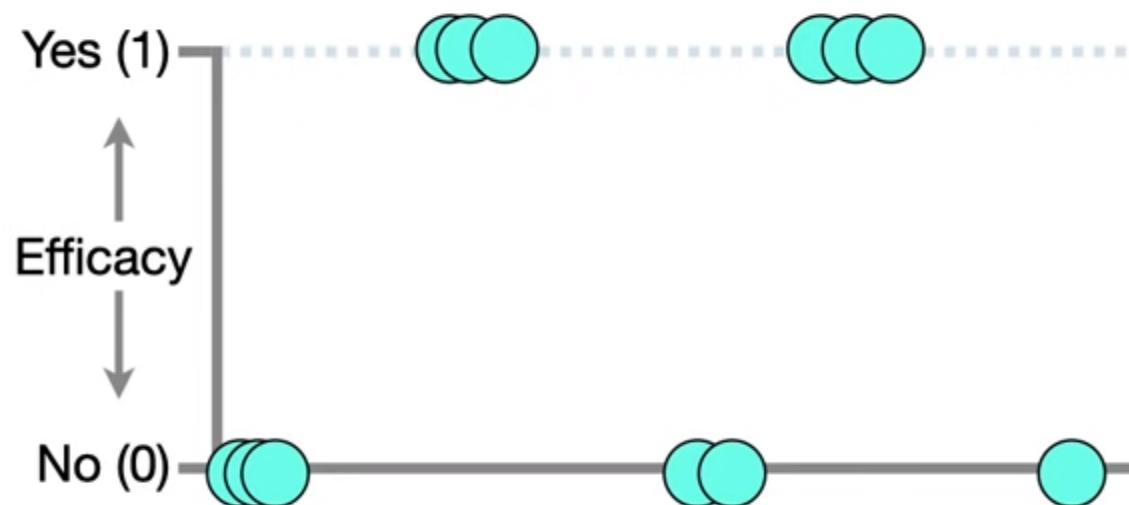
Will you go to the movies? **No**



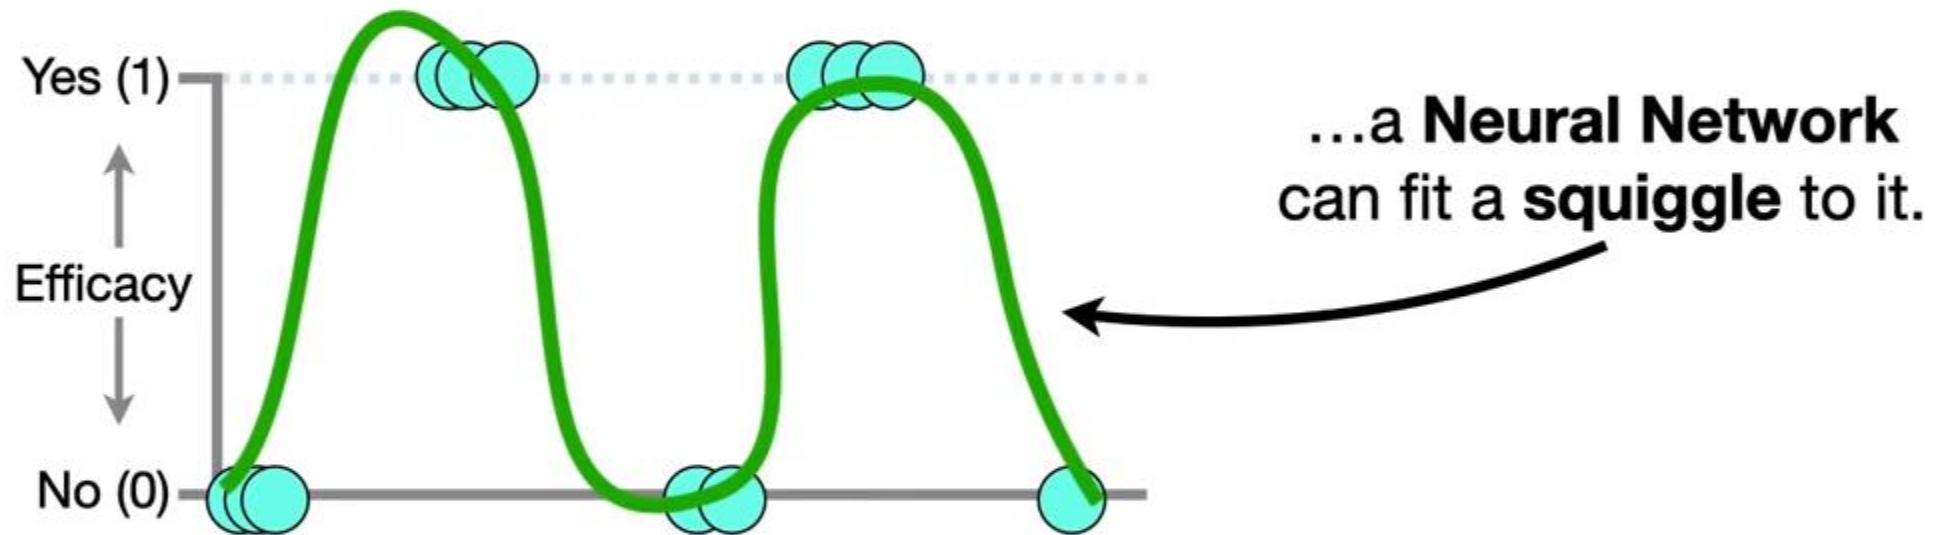


Now that we have this data, we would like to use it to predict whether or not a future **Dosage** will be **Effective**.



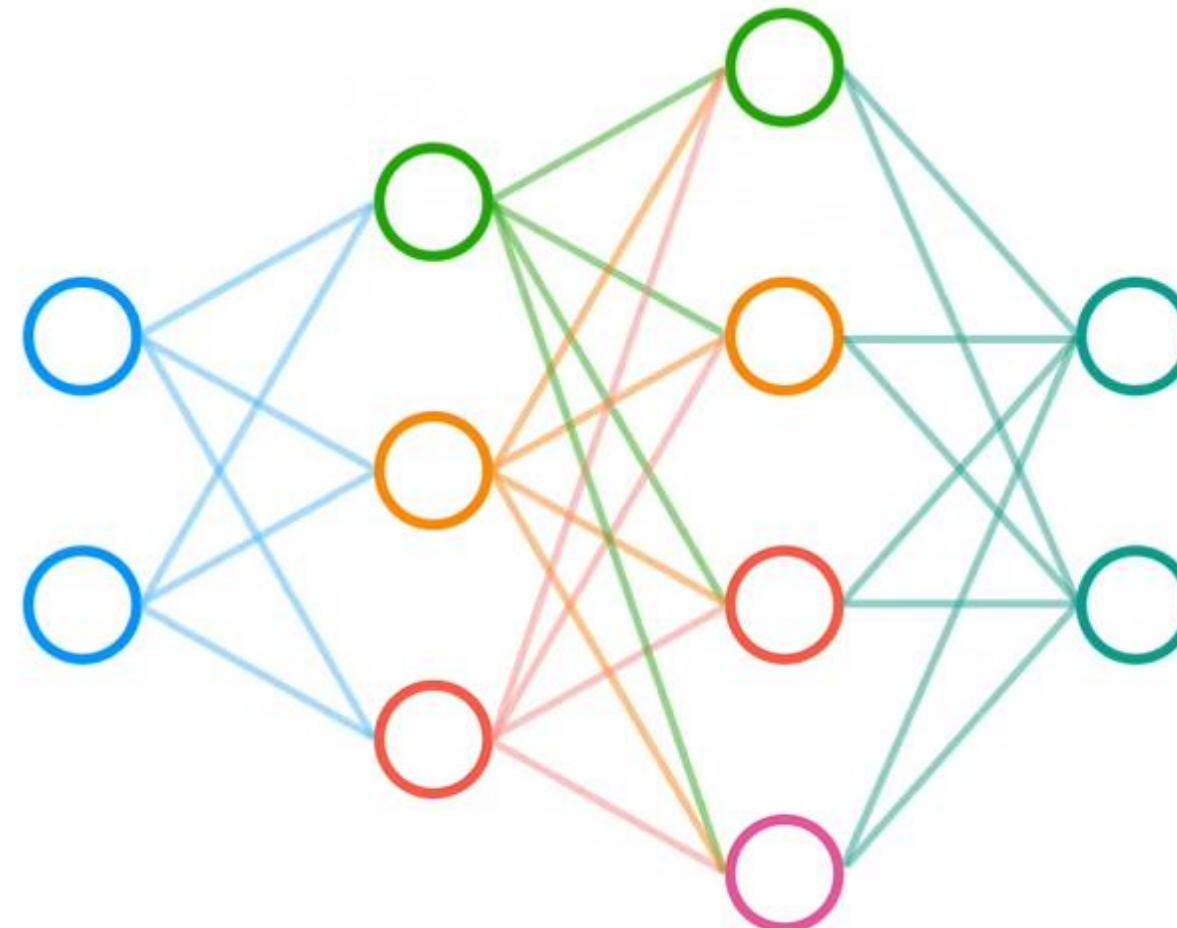


And even if we have
really complicated
dataset like this...



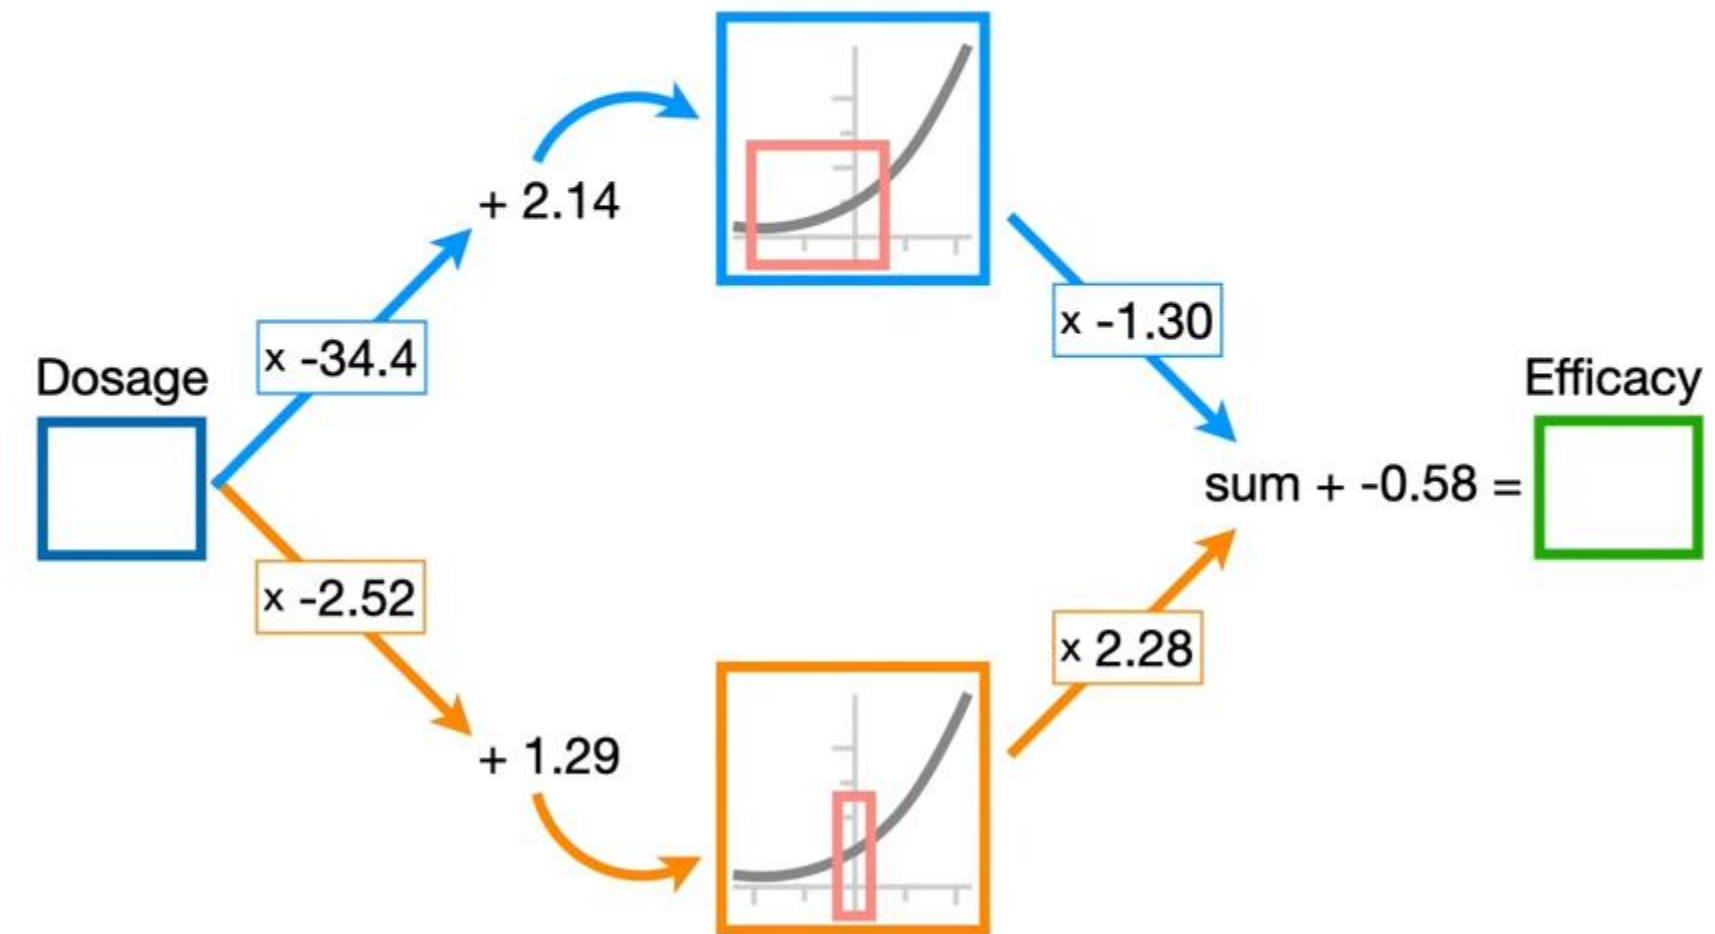
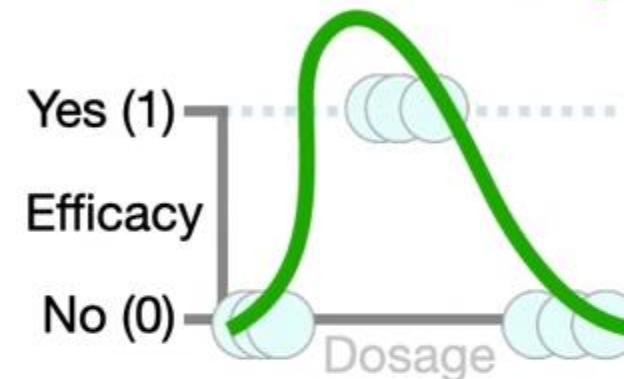


Now, even though this **Neural Network** looks fancy, it is still made from the same parts...



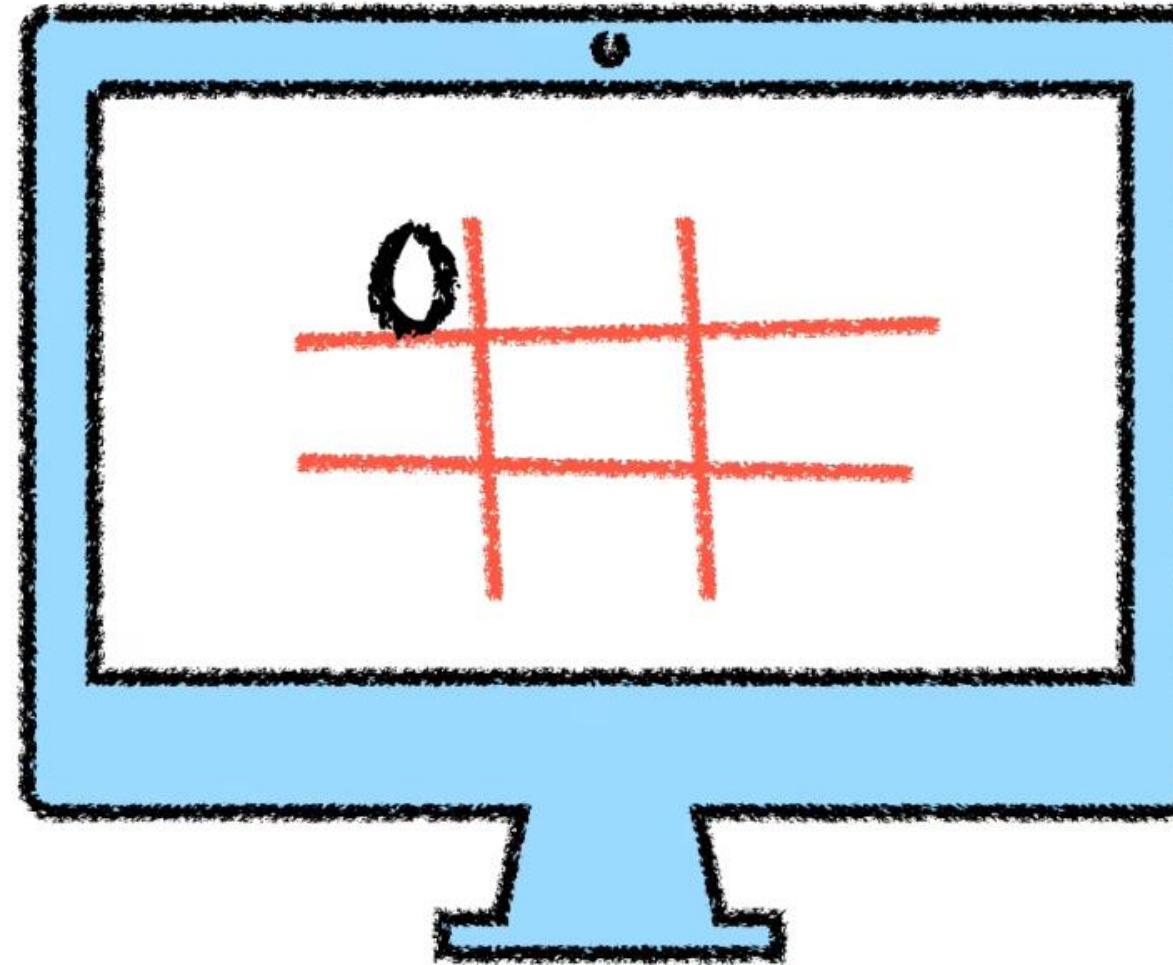


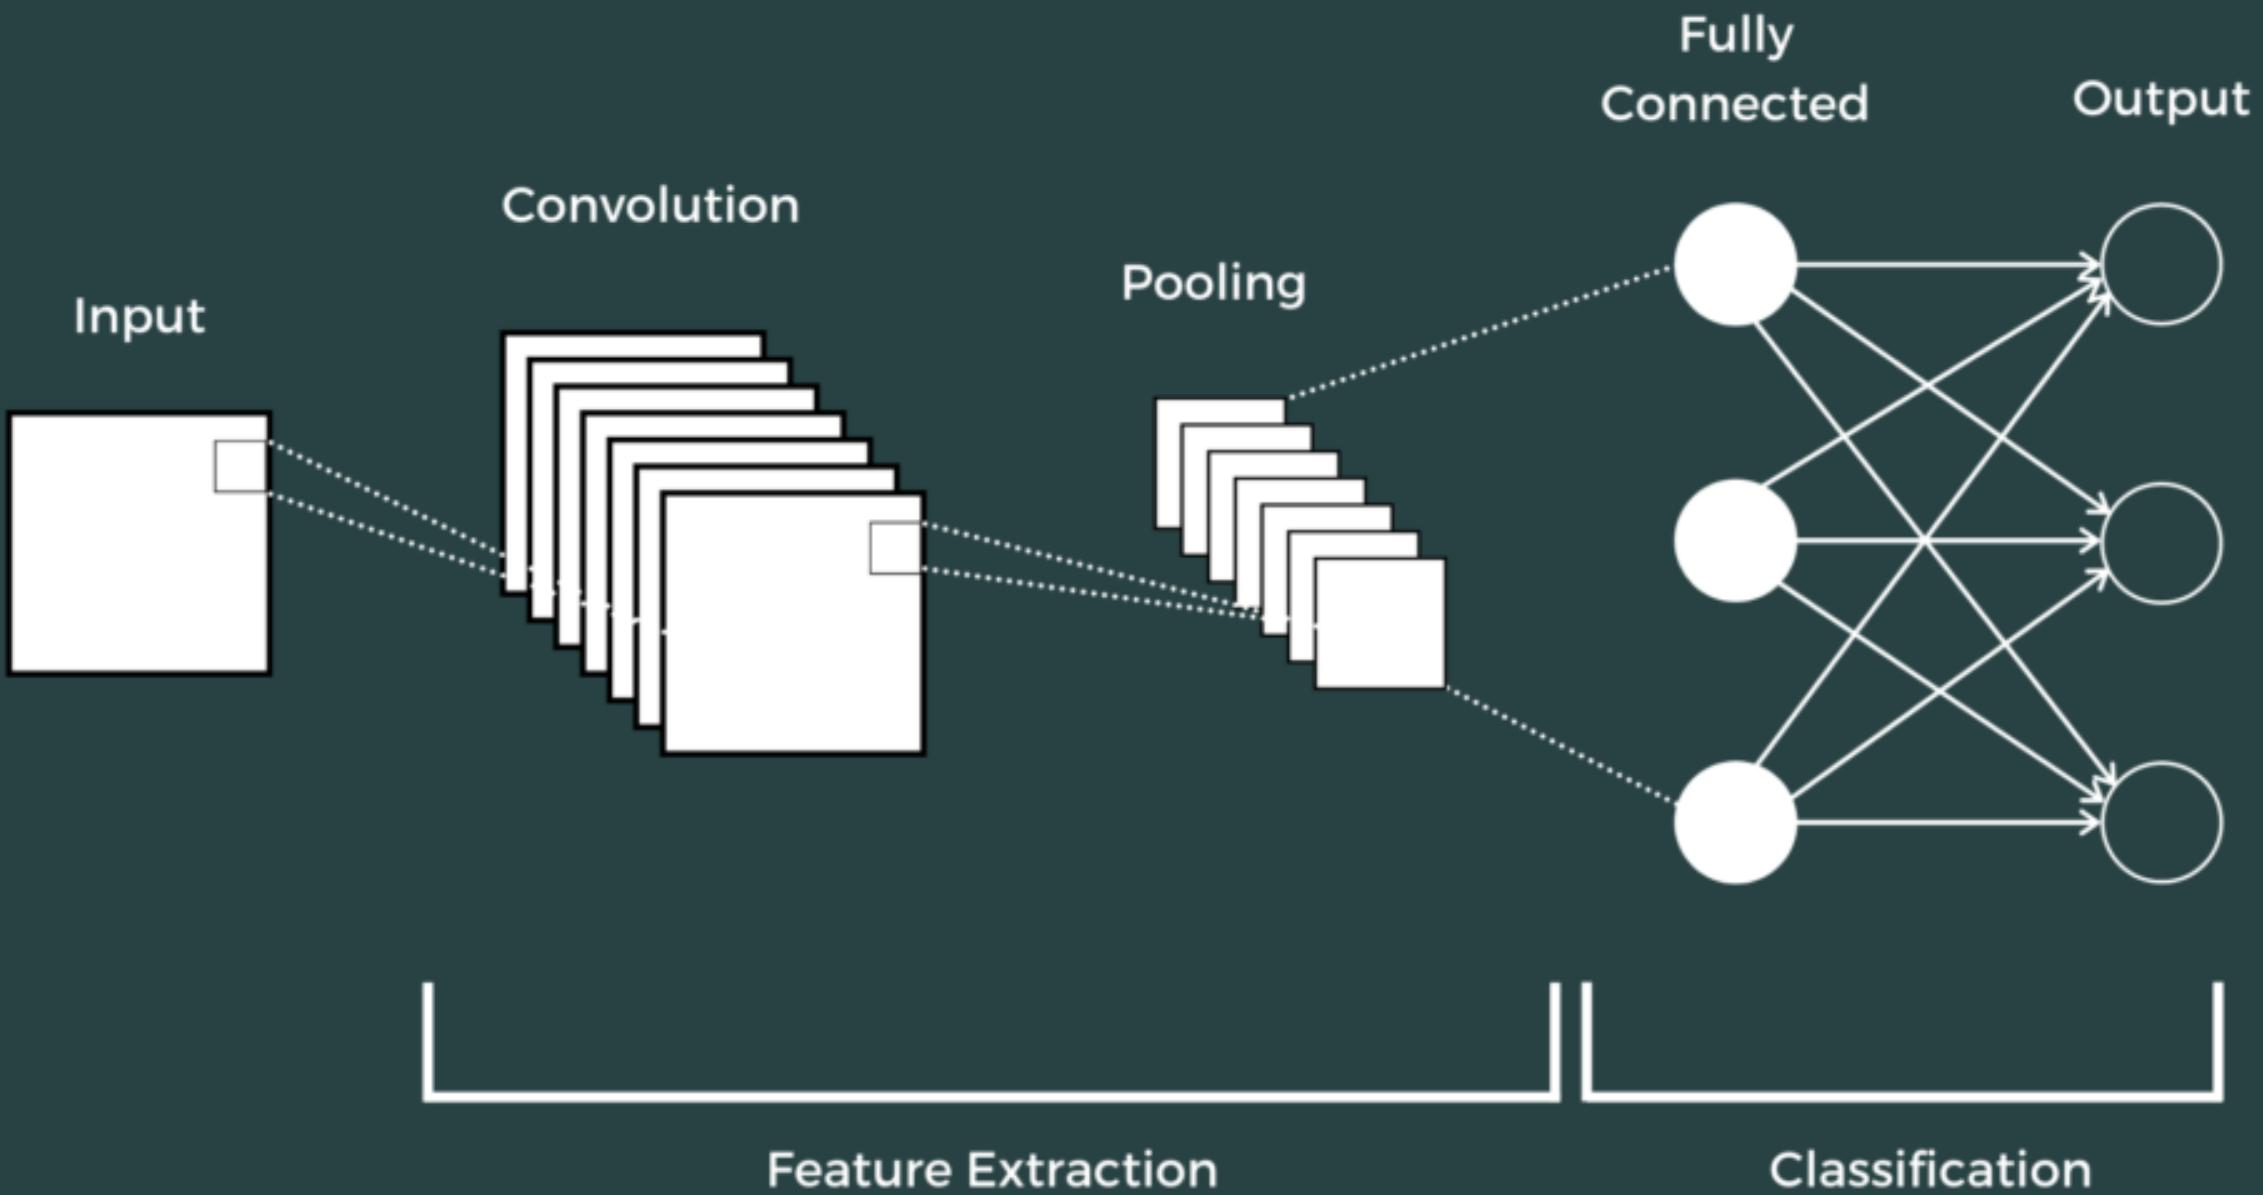
...and we have a **green** squiggle that fits the data.





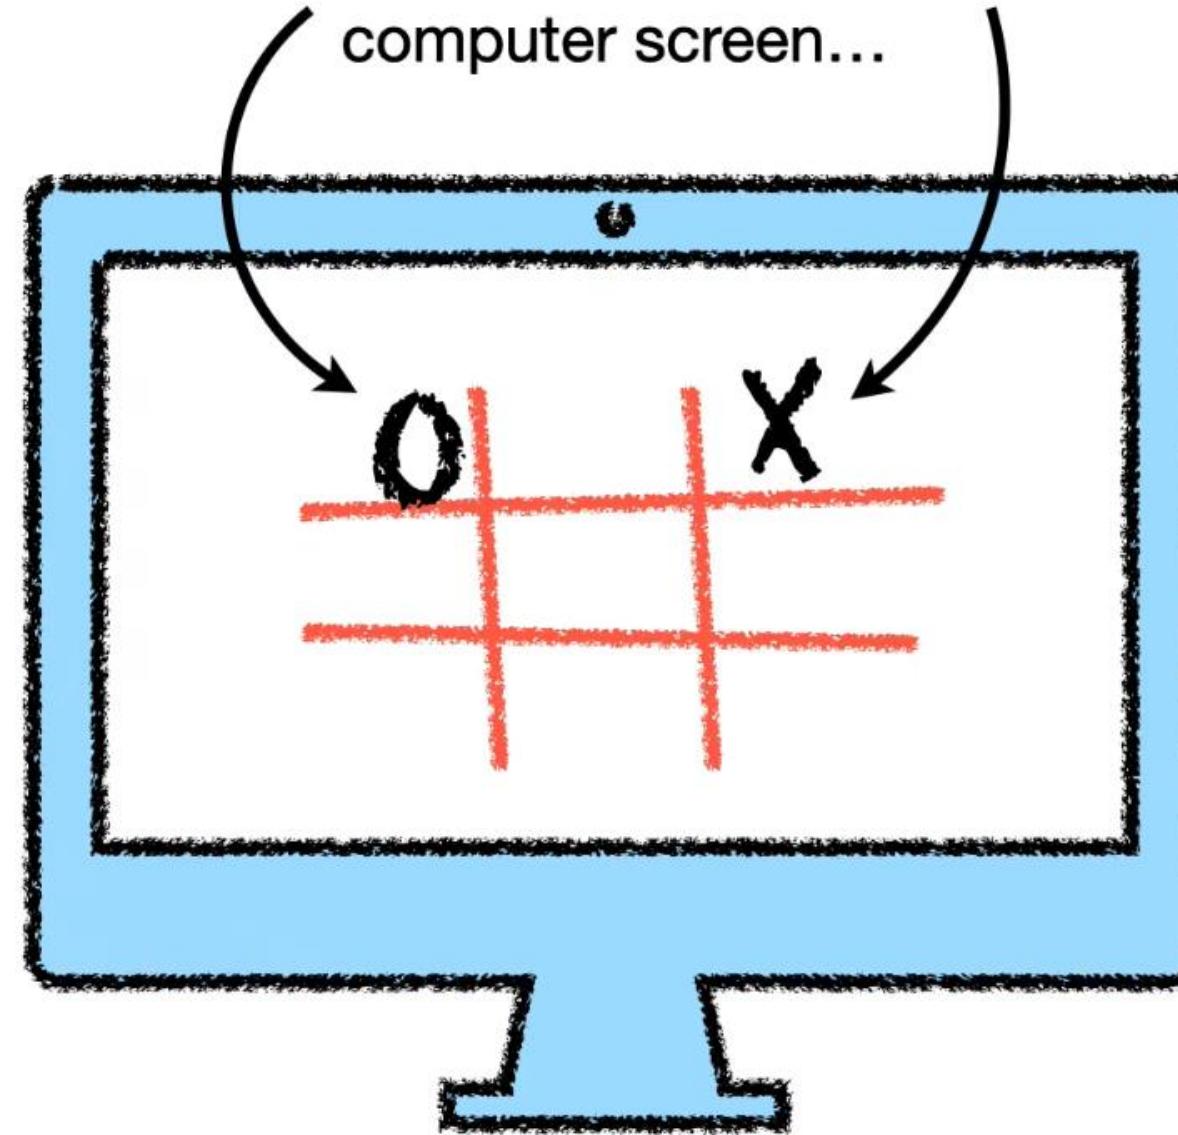
Convolutional Neural Network!!!







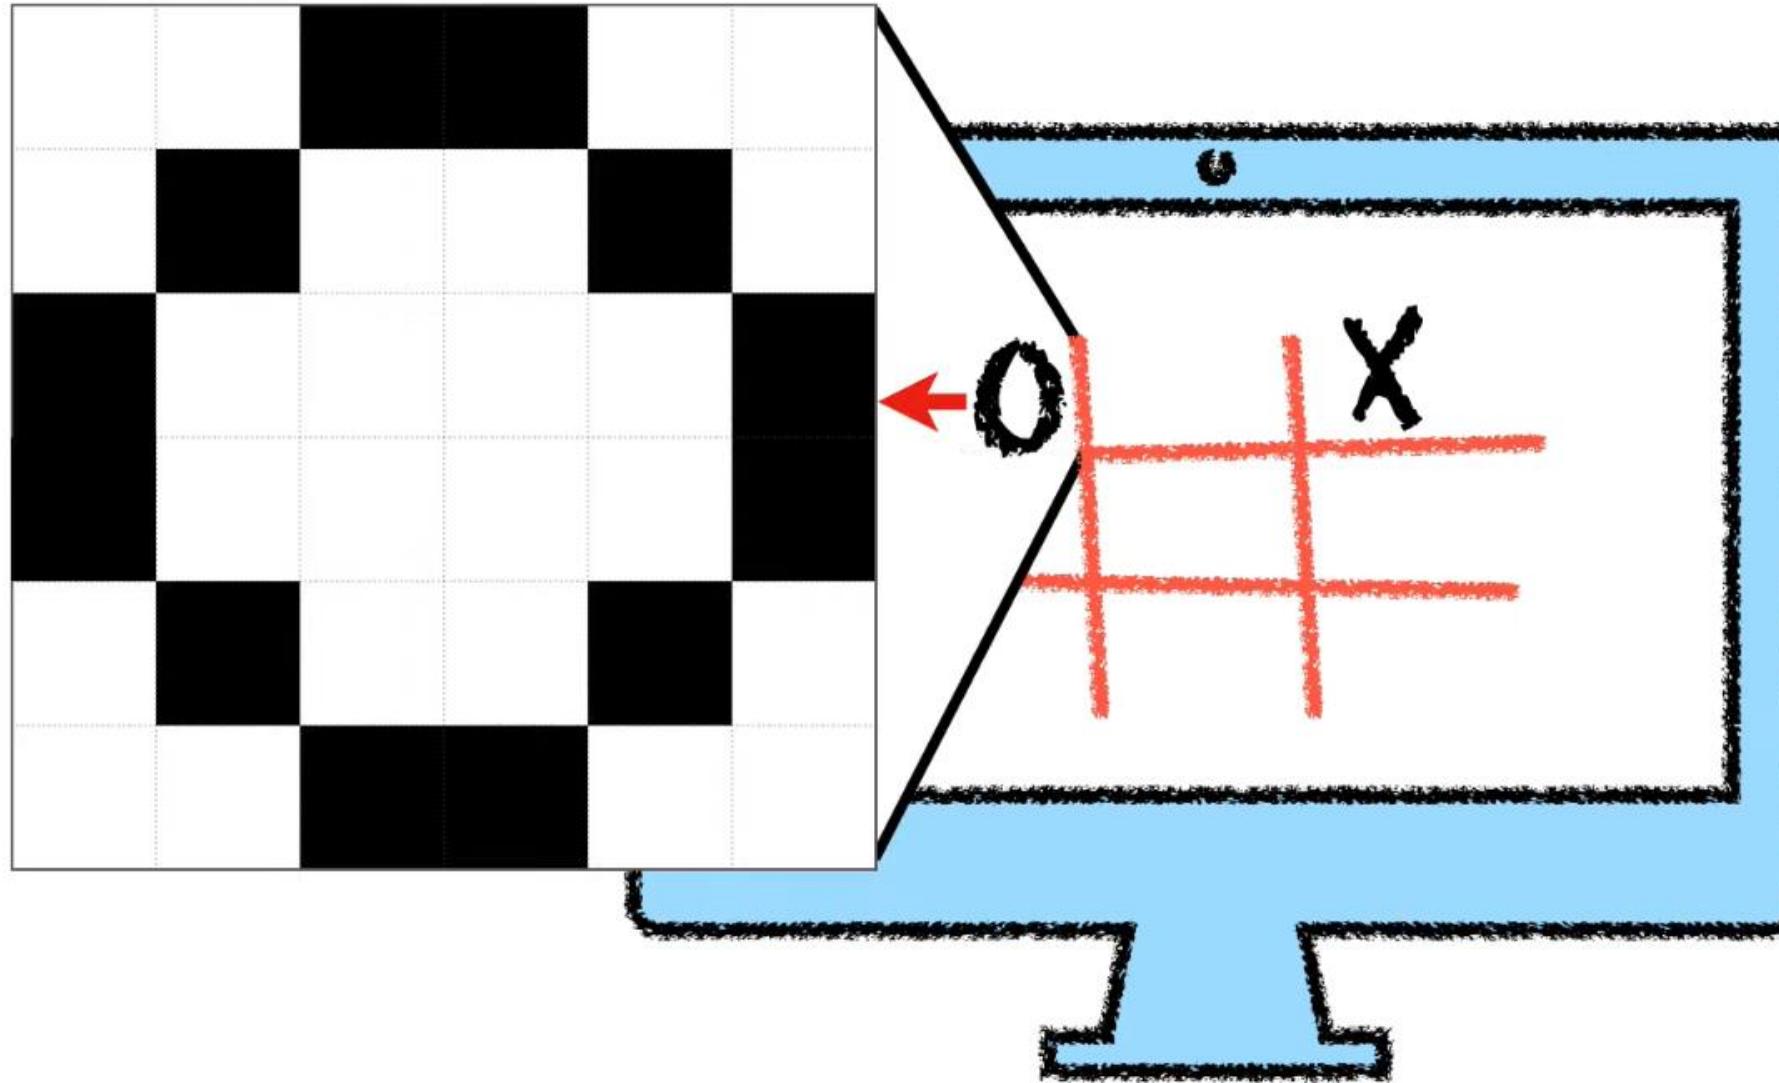
Because the letter “O” and
the letter “X” are drawn on a
computer screen...





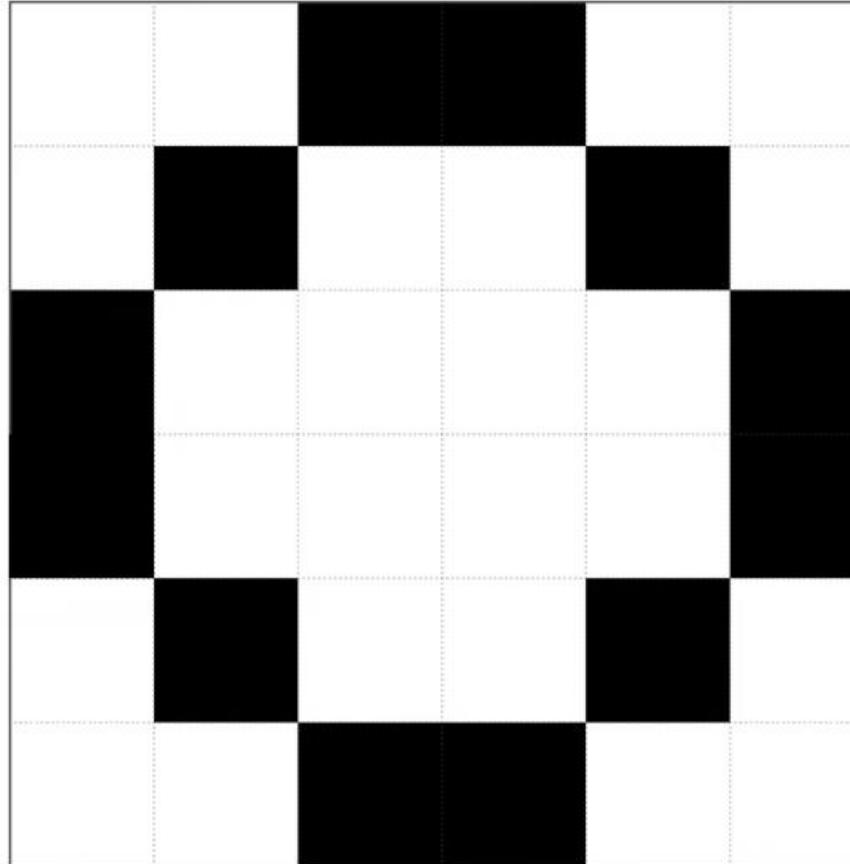
The letter “O”, zoomed in.

...we can zoom in
on the letter “O”...



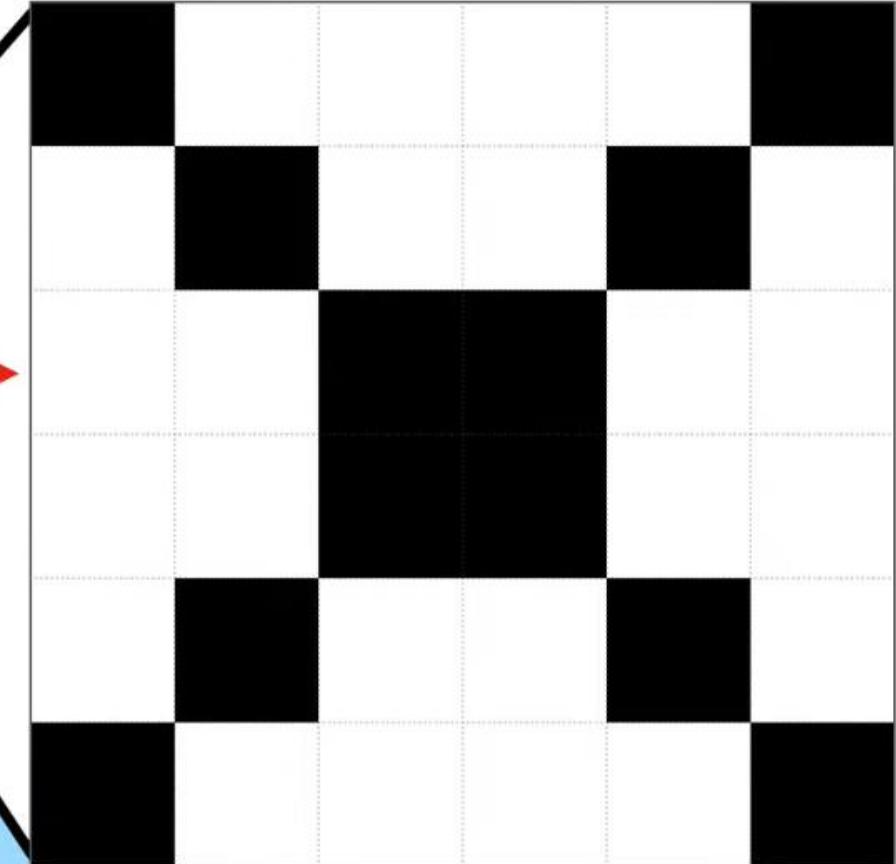
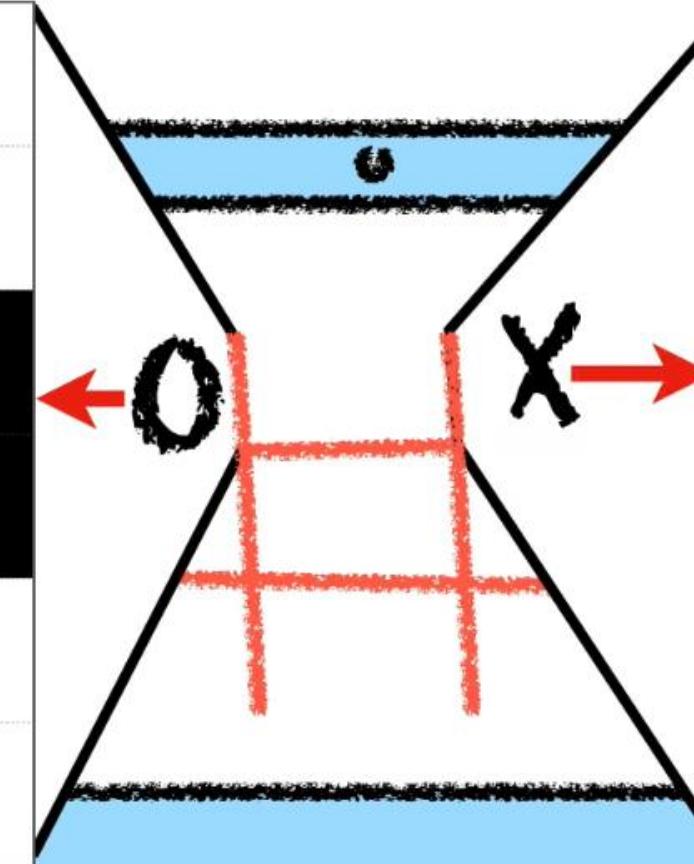


The letter “O”, zoomed in.



...and the
letter “X”...

The letter “X”, zoomed in.





The letter “O”, zoomed in.

0	0			0	0
0		0	0		0
	0	0	0	0	
	0	0	0	0	
0		0	0		0
0	0			0	0

...and, in this case, each pixel is represented by either a **0**, for a white pixel...

The letter “X”, zoomed in.

	0	0	0	0	
0		0	0	0	0
0	0		0	0	0
0	0	0		0	0
0		0	0	0	0
0	0			0	0



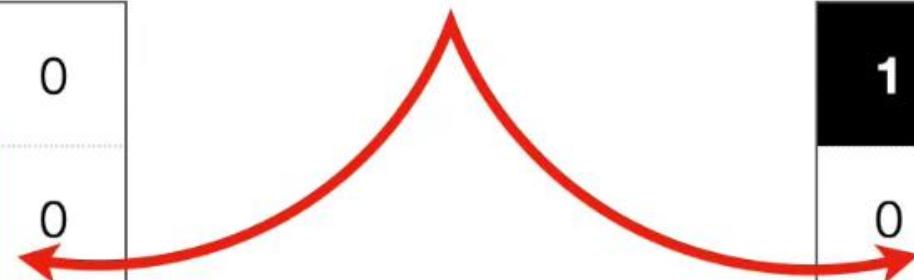
The letter “O”, zoomed in.

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

...or a **1**, for a
black pixel.

The letter “X”, zoomed in.

1	0	0	0	0	1
0	1	0	0	0	1
0	0	1	1	0	0
0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1

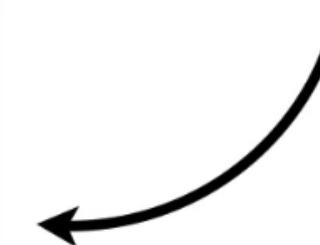




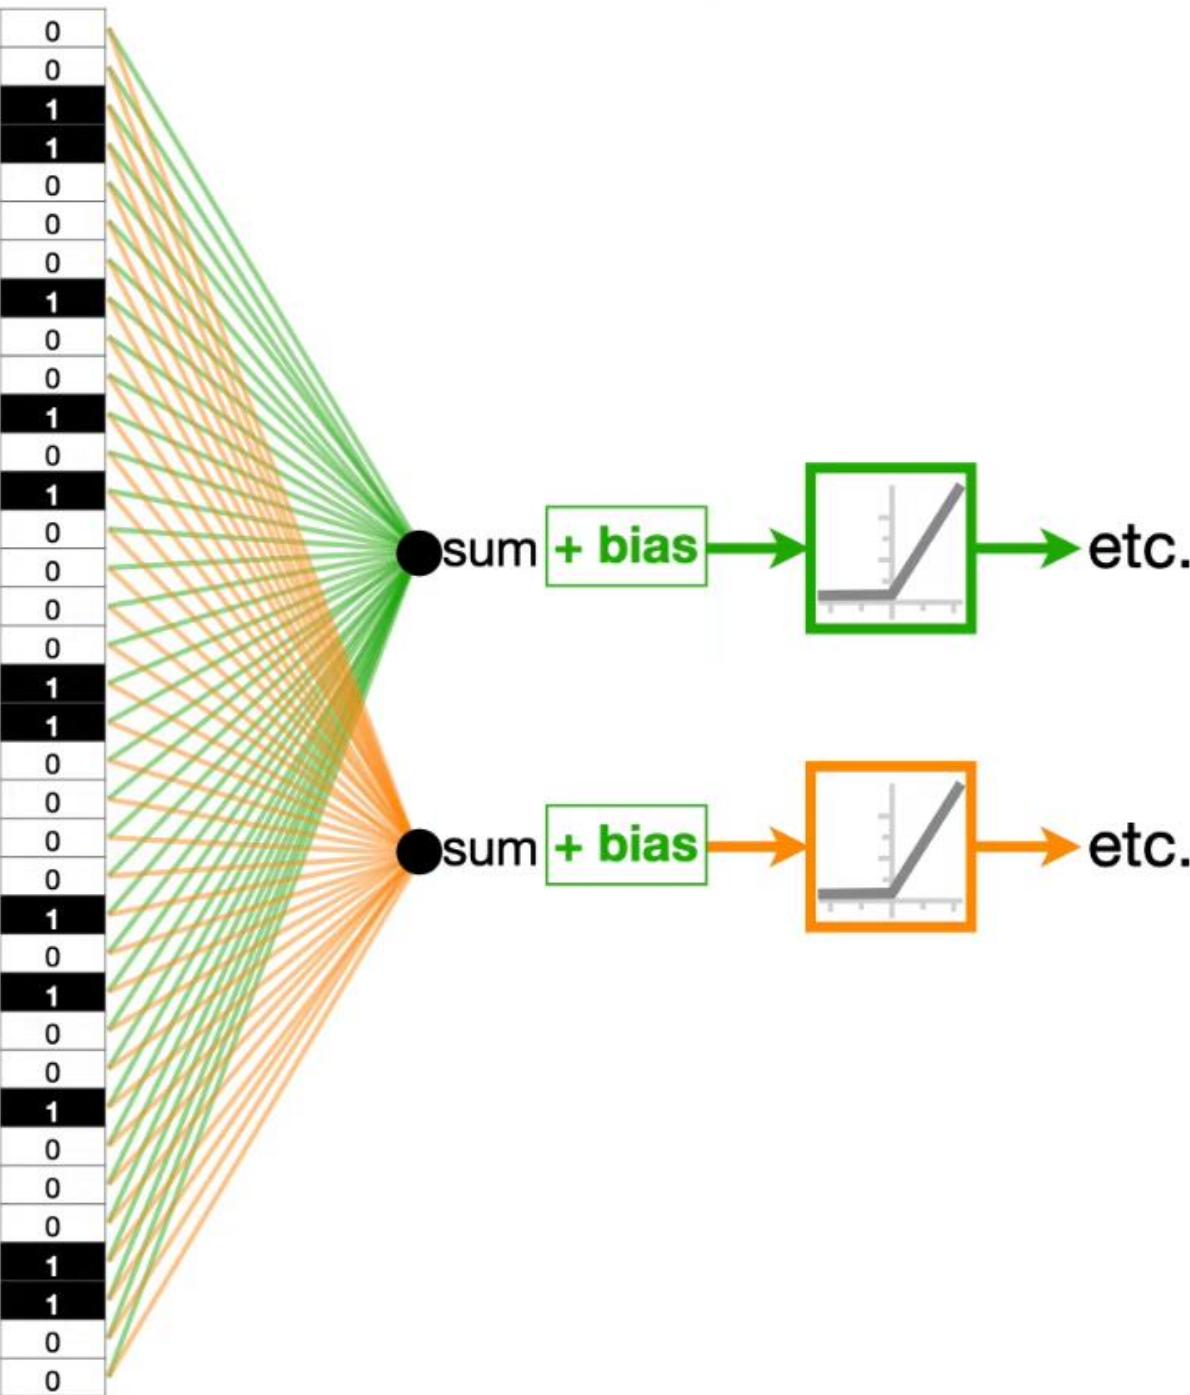
The letter “O”, zoomed in.

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

We will start with the image of the letter “O”.

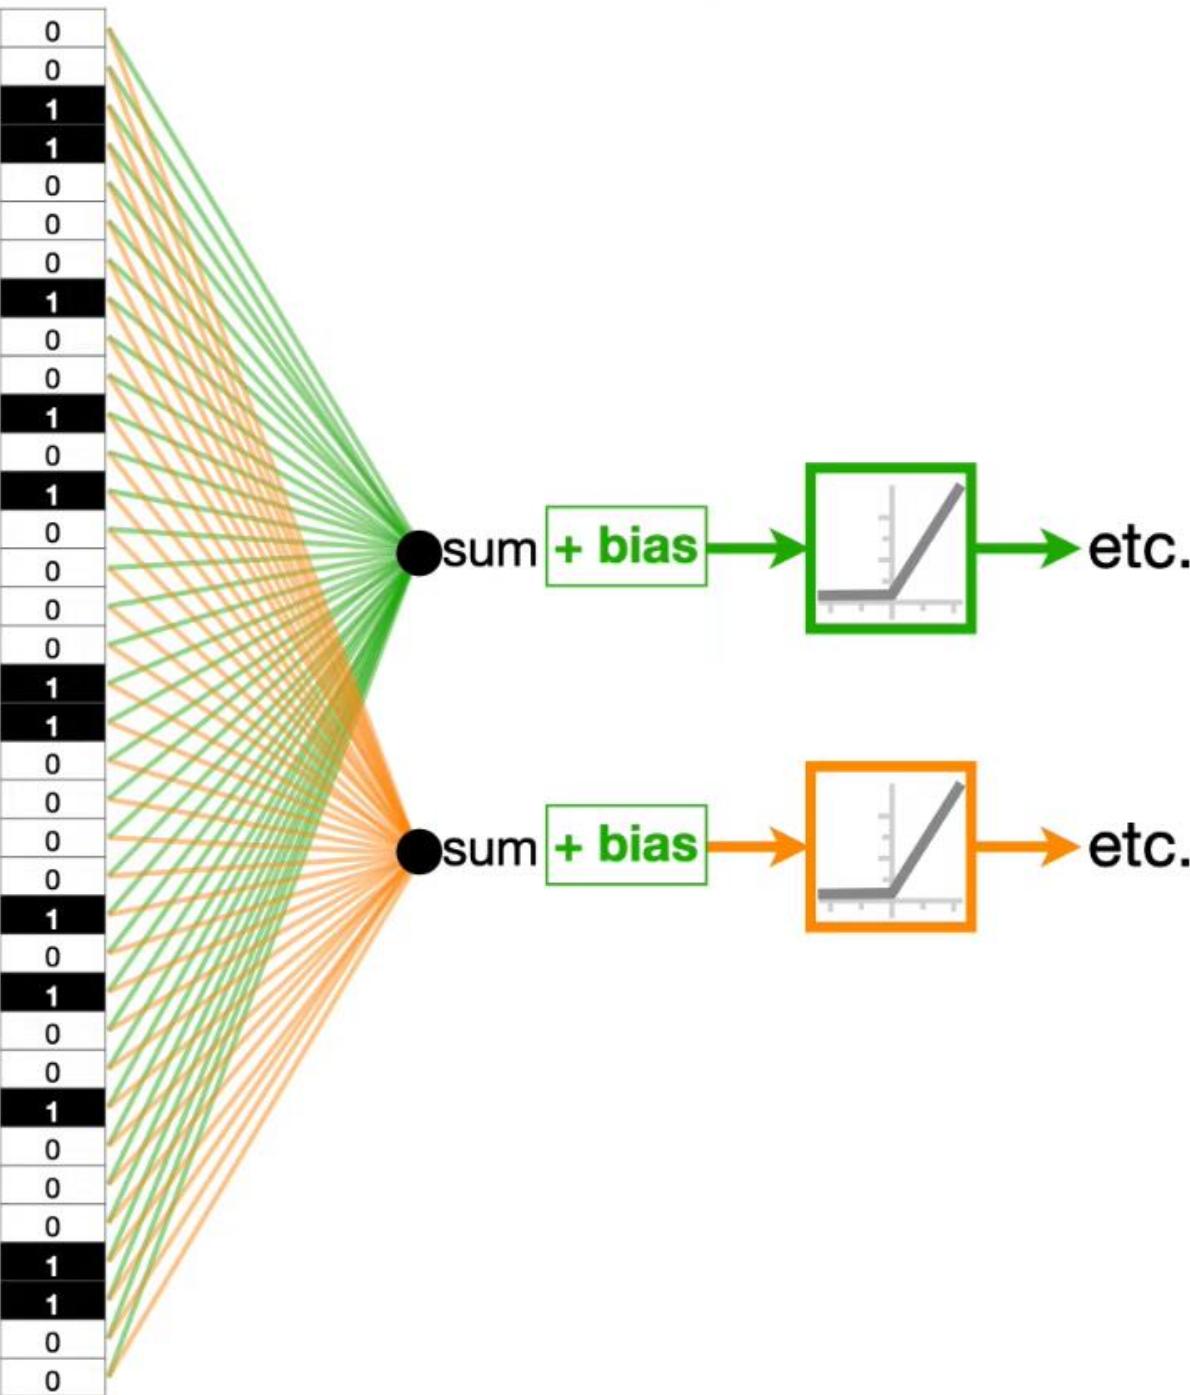


0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0



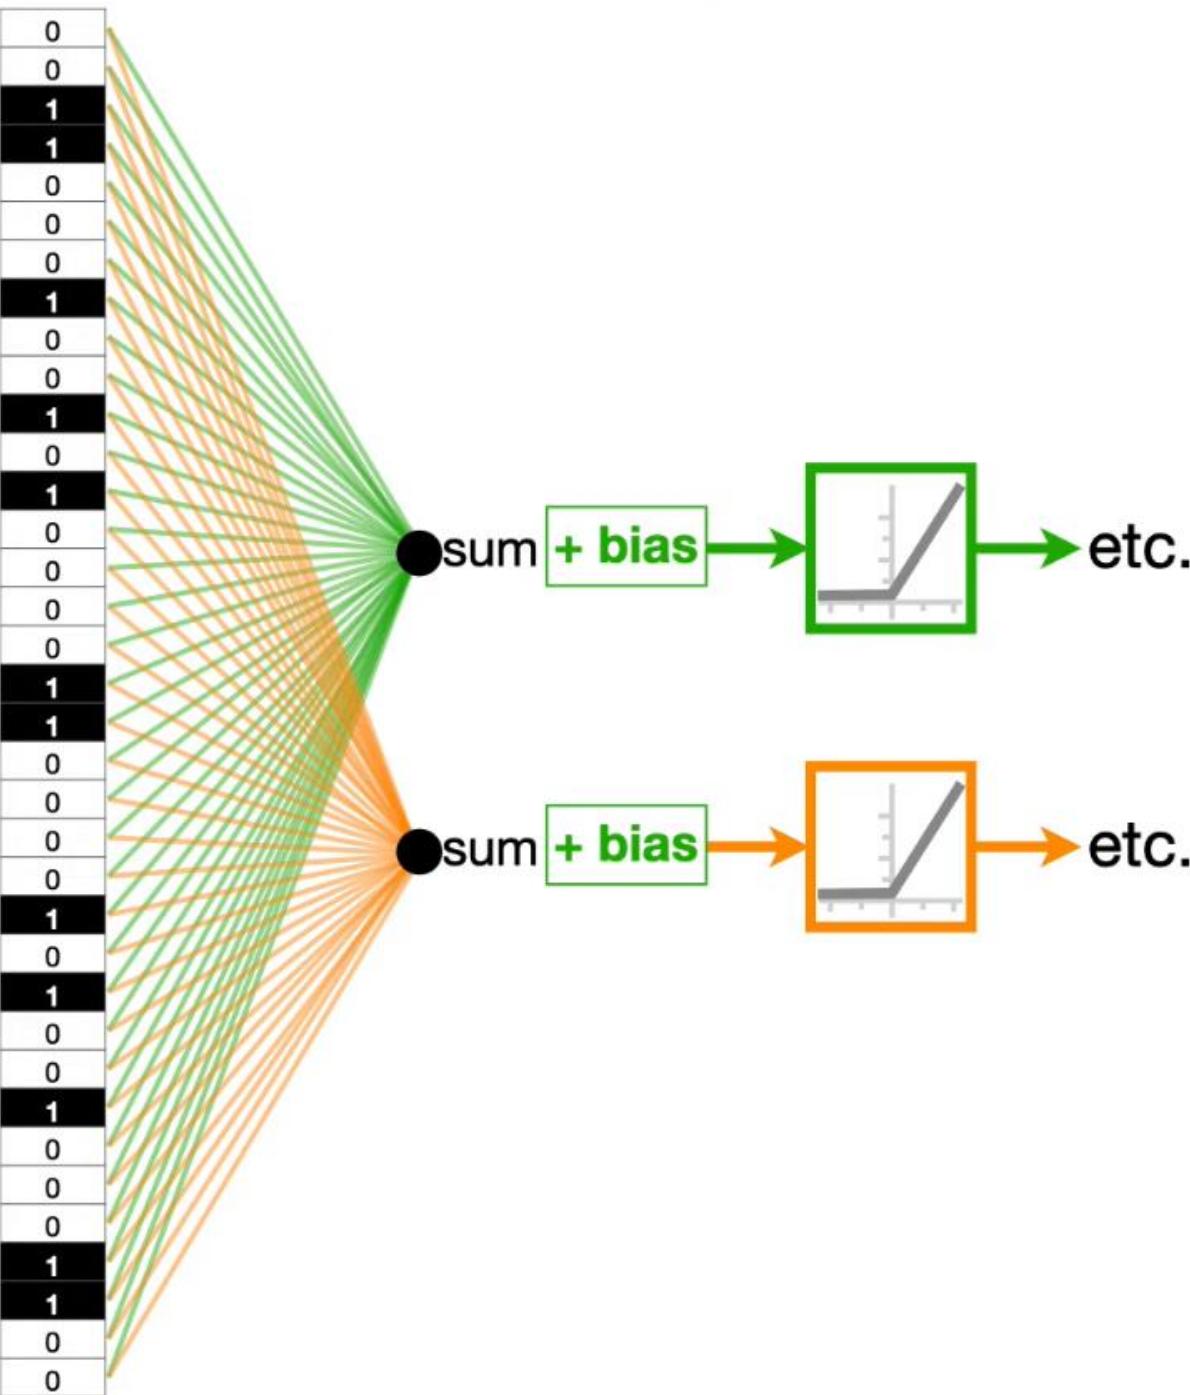
However, if we had a larger image, like **100** pixels by **100** pixels, which is still pretty small compared to real world pictures...

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0



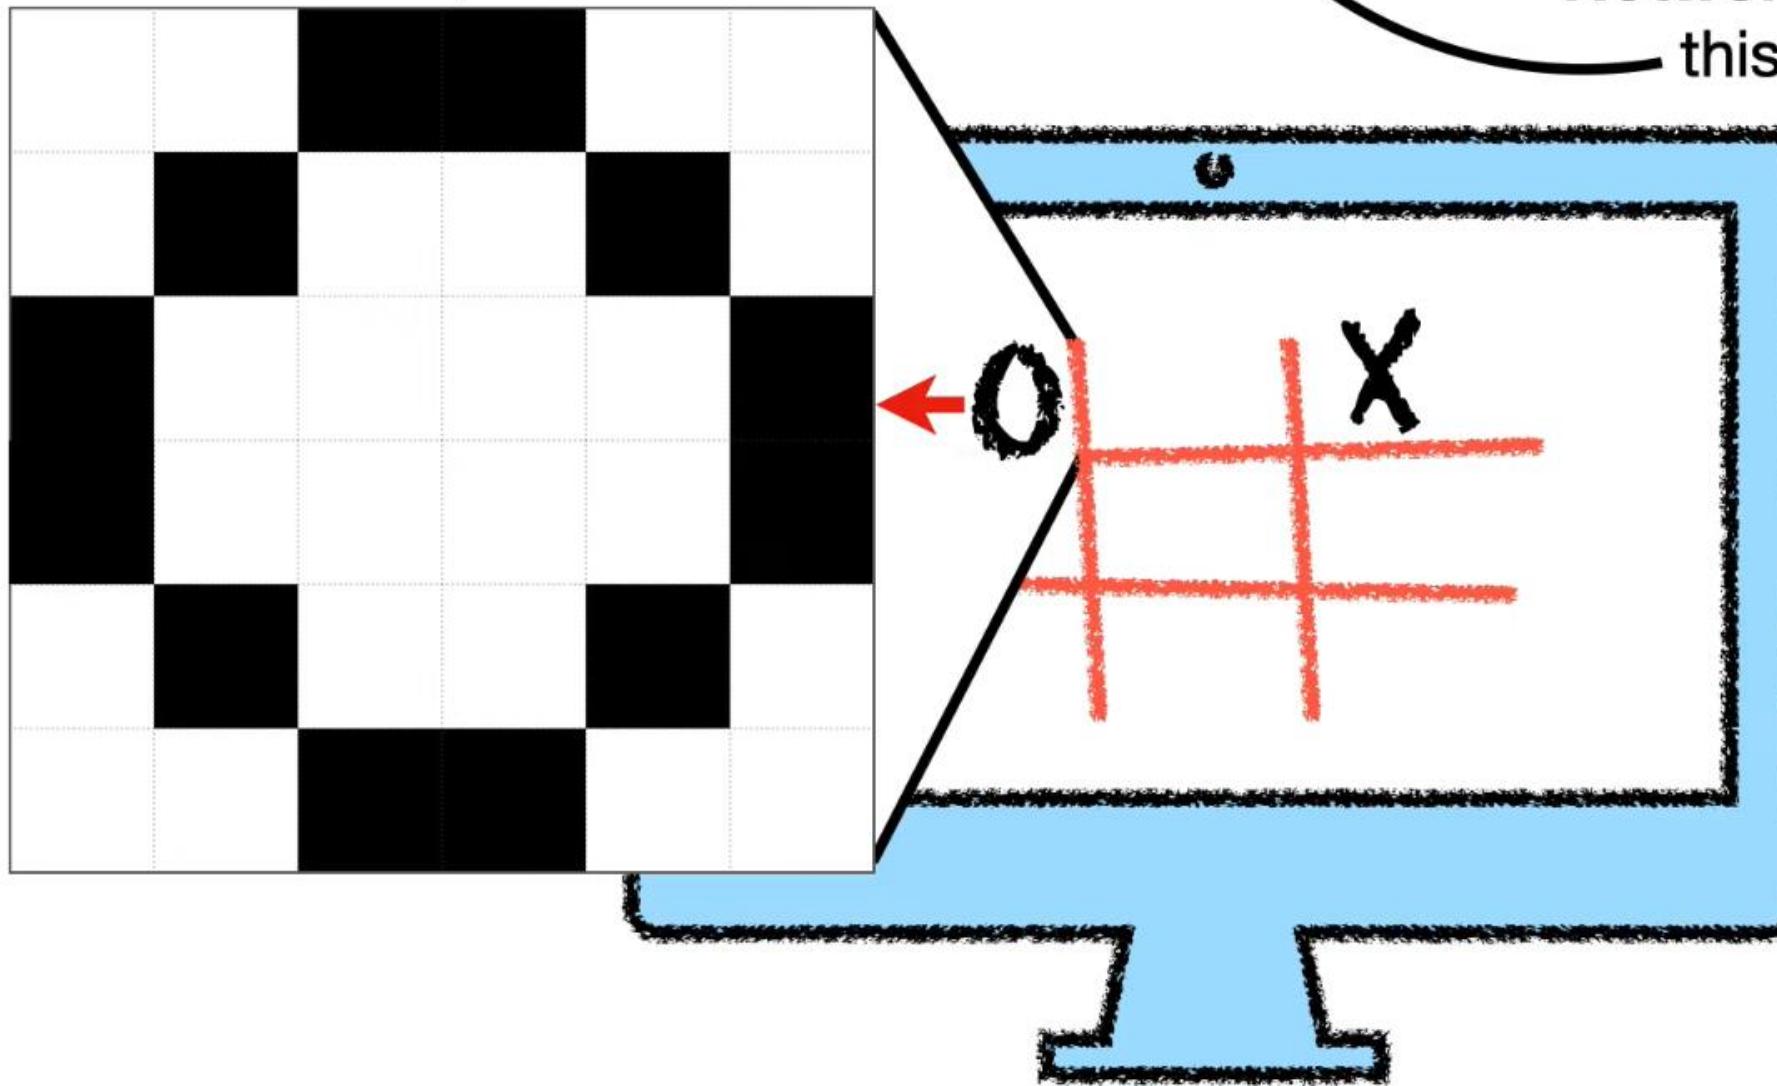
...then we would end up with having to estimate **10,000 Weights per node in the Hidden Layer!!!**

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0



So, this method
doesn't scale
very well.

The letter “O”, zoomed in.



...and see how a
**Convolutional Neural
Network** can recognize
this letter “O”!!!



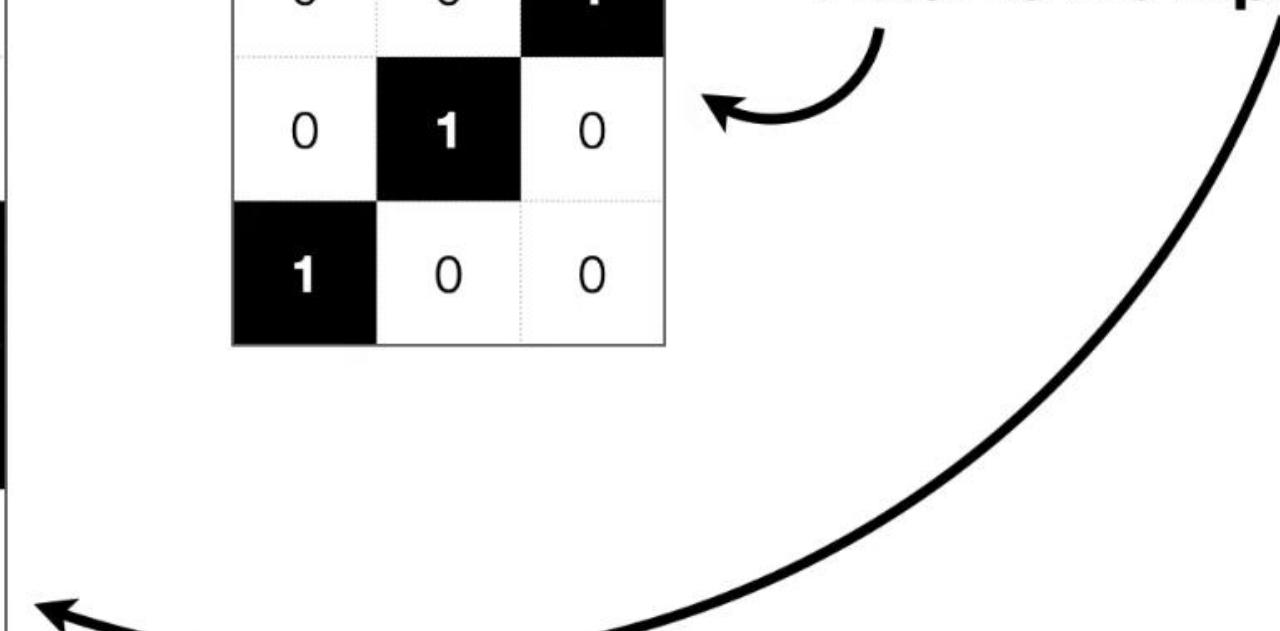
Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1
0	1	0
1	0	0

The first thing a **Convolutional Neural Network** does is apply a **Filter** to the **Input Image**.





Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1	1	0	0
0	1	0	1	0	0
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

To apply the **Filter** to the input image, we overlay the **Filter** onto the image...



Input Image

0	0	1	1	0	0	0
0	1	0	0	1	0	0
1	0	0	0	0	1	1
1	0	0	0	0	1	0
0	1	0	0	1	0	0
0	0	1	1	0	0	0

Filter (aka Kernel)

0	0	1	1	0	0	0
0	1	0	0	1	0	0
1	0	0	0	0	1	1
1	0	0	0	0	1	0

...and then we multiply together each overlapping pixel...

$$(0 \times 0) \quad (0 \times 0)$$



Input Image

0	0	1				
0	1	0				
1	0	0				
1	0	0	0	0	1	
0	1	0	0	1	0	
0	0	1	1	0	0	

Filter (aka Kernel)

0	0	1				
0	1	0				
1	0	0				
0	0	1				
0	0	0				

...and then we multiply together each overlapping pixel...

$$(0 \times 0)$$

$$(0 \times 0)$$

$$(1 \times 1)$$



Input Image

0	0	1				
0	1	0				
1	0	0				
1	0	0	0	0	1	
0	1	0	0	1	0	
0	0	1	1	0	0	

Filter (aka Kernel)

0	0	1	
0	1	0	
1	0	0	

...and then we multiply together each overlapping pixel...

$$\begin{matrix} (0 \times 0) & (0 \times 0) & (1 \times 1) \\ (0 \times 0) & (1 \times 1) \end{matrix}$$



Input Image

0	0	1		1	0	0
0	1	0		0	1	0
1	0	0		0	0	1
1	0	0		0	0	1
0	1	0		0	1	0
0	0	1	1	0	0	0

Filter (aka Kernel)

0	0	1	
0	1	0	
1	0	0	

...and then we multiply together each overlapping pixel...

$$\begin{array}{l} (0 \times 0) \quad (0 \times 0) \quad (1 \times 1) \\ (0 \times 0) \quad (1 \times 1) \quad (0 \times 0) \end{array}$$



Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	0
0	1	0	1	0	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1
0	1	0
1	0	0

...and then we multiply together each overlapping pixel...

$$\begin{array}{c} (0 \times 0) \quad (0 \times 0) \quad (1 \times 1) \\ (0 \times 0) \quad (1 \times 1) \quad (0 \times 0) \\ (1 \times 1) \end{array}$$



Input Image

0	0	1		1	0	0
0	1	0		0	1	0
1	0	0		0	0	1
1	0	0		0	0	1
0	1	0		0	1	0
0	0	1		1	0	0

Filter (aka Kernel)

0	0	1	
0	1	0	
1	0	0	

...and then we multiply together each overlapping pixel...

$$\begin{array}{c} (0 \times 0) \quad (0 \times 0) \quad (1 \times 1) \\ (0 \times 0) \quad (1 \times 1) \quad (0 \times 0) \\ \hline (1 \times 1) \quad (0 \times 0) \end{array}$$



Input Image

0	0	1		1	0	0
0	1	0		0	1	0
1	0	0		0	0	1
1	0	0		0	0	1
0	1	0		0	1	0
0	0	1		1	0	0

Filter (aka Kernel)

0	0	1
0	1	0
1	0	0

...and then we multiply together each overlapping pixel...

$$\begin{array}{l} (0 \times 0) \quad (0 \times 0) \quad (1 \times 1) \\ (0 \times 0) \quad (1 \times 1) \quad (0 \times 0) \\ (1 \times 1) \quad (0 \times 0) \quad (0 \times 0) \end{array}$$



Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1
0	1	0
1	0	0

...and then we add each product together...

$$\begin{aligned} & (0 \times 0) + (0 \times 0) + (1 \times 1) \\ & + (0 \times 0) + (1 \times 1) + (0 \times 0) \\ & + (1 \times 1) + (0 \times 0) + (0 \times 0) \end{aligned}$$



Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1
0	1	0
1	0	0

...to get a final value,
which in this case, is **3**.

$$\begin{aligned} & (0 \times 0) + (0 \times 0) + (1 \times 1) \\ & + (0 \times 0) + (1 \times 1) + (0 \times 0) \\ & + (1 \times 1) + (0 \times 0) + (0 \times 0) \end{aligned}$$

$$= 3$$



Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1
0	1	0
1	0	0

Terminology Alert!!!

In fancy math lingo, we call this sum of products a **Dot Product**.

$$\begin{aligned} & (0 \times 0) + (0 \times 0) + (1 \times 1) \\ & + (0 \times 0) + (1 \times 1) + (0 \times 0) \\ & + (1 \times 1) + (0 \times 0) + (0 \times 0) \\ & = 3 \end{aligned}$$



Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1
0	1	0
1	0	0

By computing the **Dot Product** between the input and the **Filter**, we can say that the **Filter** is **Convolved** with the input, and that's what gives **Convolutional Neural Networks** their name.

$$\begin{aligned} & (0 \times 0) + (0 \times 0) + (1 \times 1) \\ & + (0 \times 0) + (1 \times 1) + (0 \times 0) \\ & + (1 \times 1) + (0 \times 0) + (0 \times 0) \\ & = 3 \end{aligned}$$



Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1	
0	1	0	
1	0	0	

Now we add a **Bias** term
to the output of the
Filter...

$$\begin{aligned} & (0 \times 0) + (0 \times 0) + (1 \times 1) \\ & + (0 \times 0) + (1 \times 1) + (0 \times 0) \\ & + (1 \times 1) + (0 \times 0) + (0 \times 0) \end{aligned}$$

$$= 3$$



Input Image

0	0	1		1	0	0
0	1	0		0	1	0
1	0	0		0	0	1
1	0	0		0	0	1
0	1	0		0	1	0
0	0	1	1	0	0	0

Filter (aka Kernel)

0	0	1	
0	1	0	
1	0	0	

...and put the final value into something called a **Feature Map**.

Feature Map

$$\begin{aligned} & (0 \times 0) + (0 \times 0) + (1 \times 1) \\ & + (0 \times 0) + (1 \times 1) + (0 \times 0) \\ & + (1 \times 1) + (0 \times 0) + (0 \times 0) \end{aligned}$$

$$= 3$$

1						



Input Image

0	1	1						
1	0	0						
1	0	0						
1	0	0	0					
0	1	0	0	1	0			
0	0	1	1	0	0			

Filter (aka Kernel)

0	0	1						
0	1	0						
1	0	0						
1	0	0	0					
0	1	0	0	1	0			
0	0	1	1	0	0			

Now, in this example,
we slide the **Filter**
over one pixel...

+ -2

Feature Map

1								



Input Image

0	1	1		0	0
1	0	0		1	0
1	0	0		0	1
1	0	0		0	1
0	1	0		1	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1	
0	1	0	
1	0	0	

...and calculate the **Dot Product** of the **Filter** and overlapping pixels in the image...

$$\begin{aligned} & (0 \times 0) + (1 \times 0) + (1 \times 1) \\ & + (1 \times 0) + (0 \times 1) + (0 \times 0) \\ & + (0 \times 1) + (0 \times 0) + (0 \times 0) \end{aligned}$$

+ -2

Feature Map

1			



Input Image

0	1	1		0	0
1	0	0	0	1	0
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1	
0	1	0	0
1	0	0	0

...add the **Bias** term...

$$+ -2$$

Feature Map

$$\begin{aligned} & (0 \times 0) + (1 \times 0) + (1 \times 1) \\ & + (1 \times 0) + (0 \times 1) + (0 \times 0) \\ & + (0 \times 1) + (0 \times 0) + (0 \times 0) \end{aligned}$$

$$= 1$$

1			



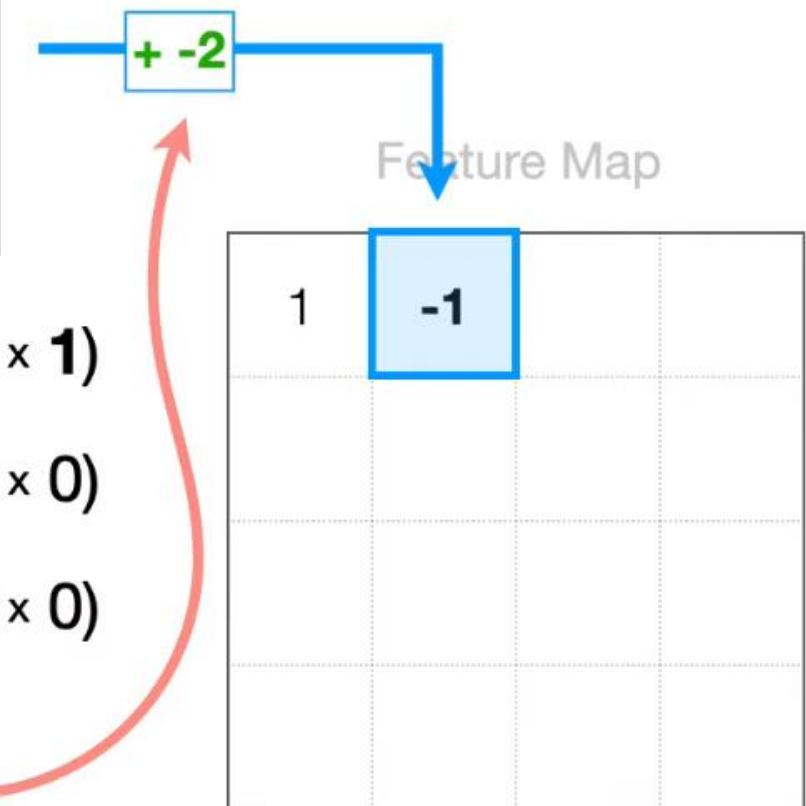
Input Image

0	1	1		0	0
1	0	0	0	1	0
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1	
0	1	0	
1	0	0	

...and put the final value into the **Feature Map**.





Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1
0	1	0
1	0	0

Then we shift the **Filter** over again and repeat until we have filled up the **Feature Map**.

$$\begin{aligned} & (1 \times 0) + (1 \times 0) + (0 \times 1) \\ & + (0 \times 0) + (0 \times 1) + (1 \times 0) \\ & + (0 \times 1) + (0 \times 0) + (0 \times 0) \end{aligned}$$

$$= 0$$



1	-1	



Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

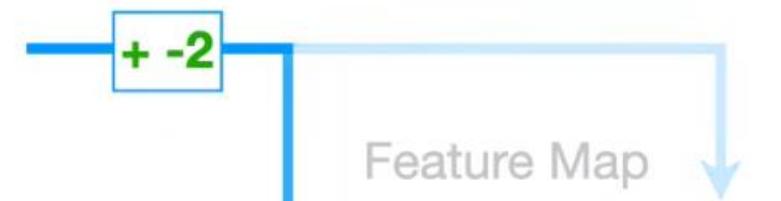
Filter (aka Kernel)

0	0	1
0	1	0
1	0	0

Then we shift the **Filter** over again and repeat until we have filled up the **Feature Map**.

$$\begin{aligned} & (0 \times 0) + (1 \times 0) + (0 \times 1) \\ & + (1 \times 0) + (0 \times 1) + (0 \times 0) \\ & + (1 \times 1) + (0 \times 0) + (0 \times 0) \end{aligned}$$

= 1



	-1	-2	-1
	-1		



Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

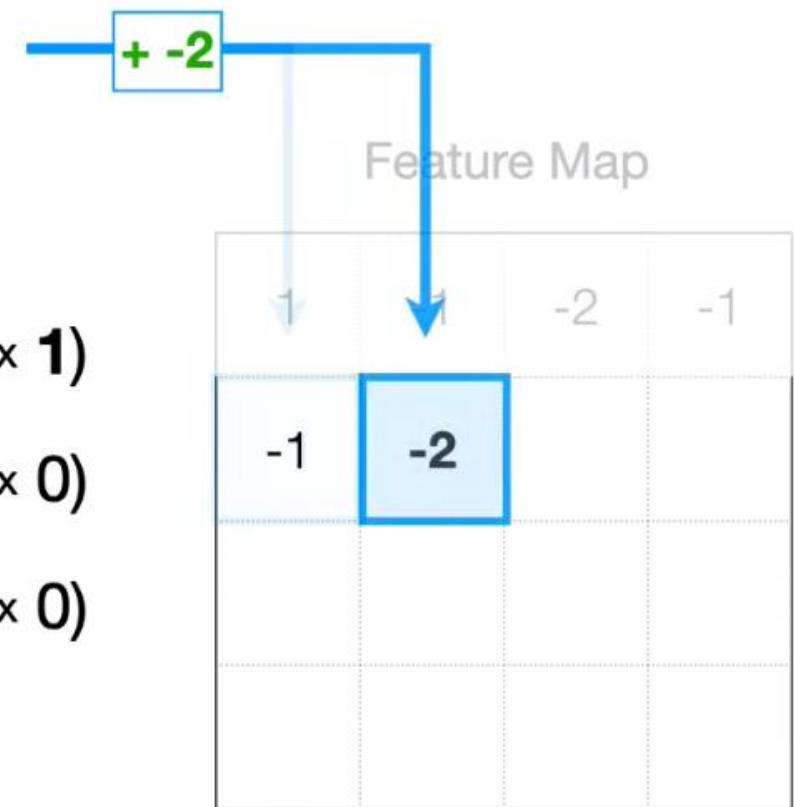
Filter (aka Kernel)

0	0	1
0	1	0
1	0	0

$$\begin{aligned} & (1 \times 0) + (0 \times 0) + (0 \times 1) \\ & + (0 \times 0) + (0 \times 1) + (0 \times 0) \\ & + (0 \times 1) + (0 \times 0) + (0 \times 0) \end{aligned}$$

$$= 0$$

Then we shift the **Filter** over again and repeat until we have filled up the **Feature Map**.





Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1
0	1	0
1	0	0

Then we shift the **Filter** over again and repeat until we have filled up the **Feature Map**.

$$\begin{aligned} & (0 \times 0) + (0 \times 0) + (1 \times 1) \\ & + (0 \times 0) + (0 \times 1) + (0 \times 0) \\ & + (0 \times 1) + (0 \times 0) + (0 \times 0) \end{aligned}$$

$$= 1$$

$$+ -2$$

Feature Map

1	-1	-2	-1
-1	-2	-1	



Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1
0	1	0
1	0	0



$$(0 \times 0) + (1 \times 0) + (0 \times 1)$$

$$+ (0 \times 0) + (0 \times 1) + (1 \times 0)$$

$$+ (0 \times 1) + (0 \times 0) + (1 \times 0)$$

$$= 0$$

Then we shift the **Filter** over again and repeat until we have filled up the **Feature Map**.

Feature Map

1	-1	-2	
-1	-2	-1	-2



Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1
0	1	0
1	0	0

Then we shift the **Filter** over again and repeat until we have filled up the **Feature Map**.

Feature Map

$$\begin{aligned} & (1 \times 0) + (0 \times 0) + (0 \times 1) \\ & + (1 \times 0) + (0 \times 1) + (0 \times 0) \\ & + (0 \times 1) + (1 \times 0) + (0 \times 0) \end{aligned}$$

$$= 0$$

-1	-2	-1	-1	-2
-2	-1	-1	-2	-1
-1	-2	-1	-1	-2
-1	-2	-1	-1	-2
-2	-1	-1	-1	-2



Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

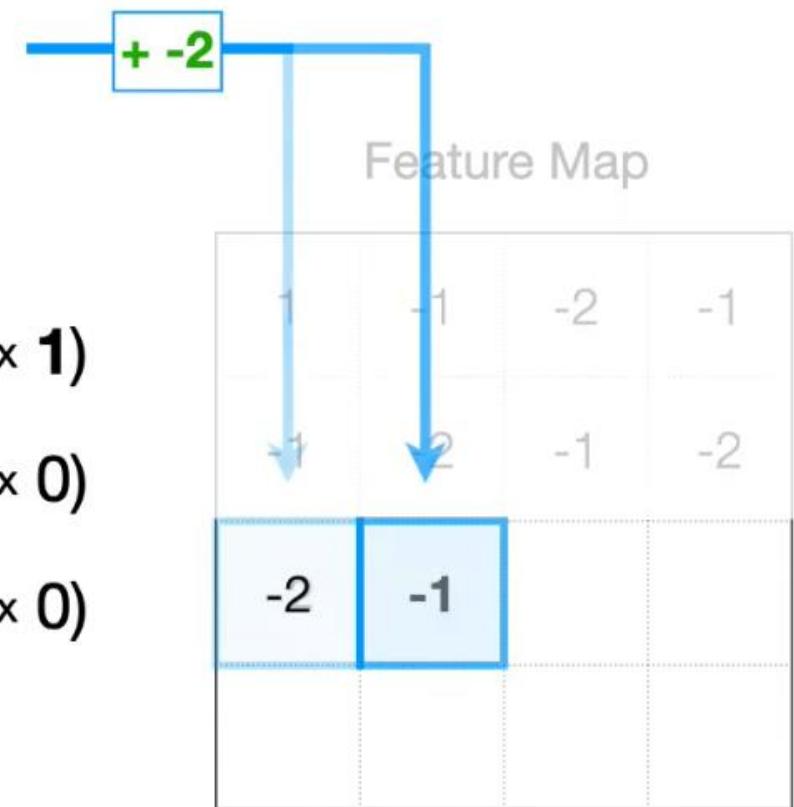
Filter (aka Kernel)

0	0	1
0	1	0
1	0	0

$$\begin{aligned} & (0 \times 0) + (0 \times 0) + (0 \times 1) \\ & + (0 \times 0) + (0 \times 1) + (0 \times 0) \\ & + (1 \times 1) + (0 \times 0) + (0 \times 0) \end{aligned}$$

$$= 0$$

Then we shift the **Filter** over again and repeat until we have filled up the **Feature Map**.





Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1
0	1	0
1	0	0

$$(0 \times 0) + (0 \times 0) + (0 \times 1)$$

$$+ (0 \times 0) + (0 \times 1) + (0 \times 0)$$

$$+ (1 \times 1) + (0 \times 0) + (0 \times 0)$$

$$= 1$$

Then we shift the **Filter** over again and repeat until we have filled up the **Feature Map**.

$$+ -2$$

Feature Map

1	-1	-2	-1
-1	2	-1	-2
-2	-1	-2	



Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

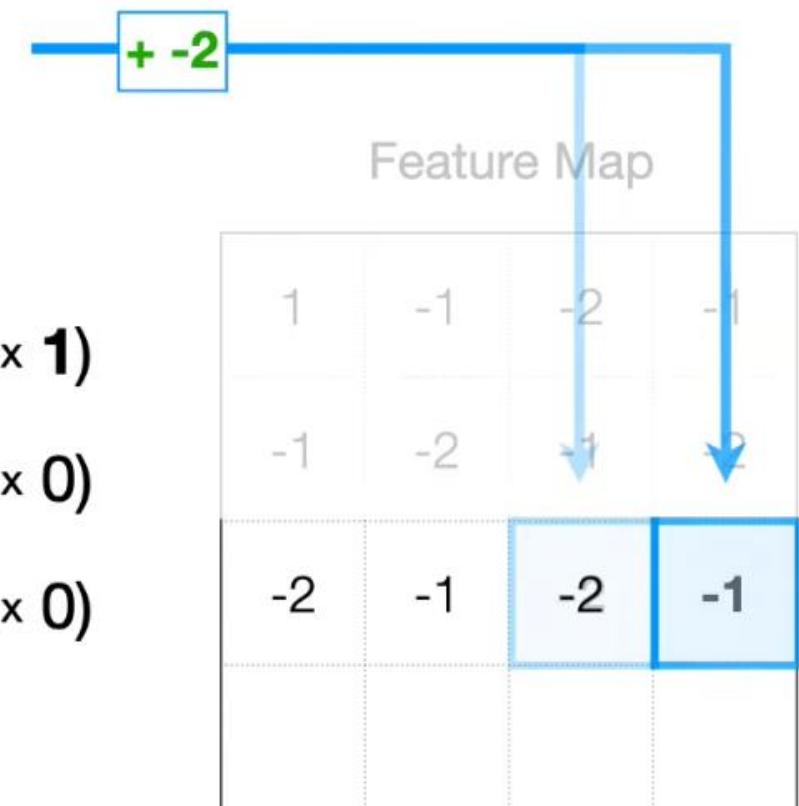
Filter (aka Kernel)

0	0	1
0	1	0
1	0	0

$$\begin{aligned} & (0 \times 0) + (0 \times 0) + (1 \times 1) \\ & + (0 \times 0) + (0 \times 1) + (1 \times 0) \\ & + (0 \times 1) + (1 \times 0) + (0 \times 0) \end{aligned}$$

$$= 1$$

Then we shift the **Filter** over again and repeat until we have filled up the **Feature Map**.





Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1
0	1	0
1	0	0

Then we shift the **Filter** over again and repeat until we have filled up the **Feature Map**.

$$\begin{aligned} & (1 \times 0) + (0 \times 0) + (0 \times 1) \\ & + (0 \times 0) + (1 \times 1) + (0 \times 0) \\ & + (0 \times 1) + (0 \times 0) + (1 \times 0) \end{aligned}$$

$$= 1$$





Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1
0	1	0
1	0	0

Then we shift the **Filter** over again and repeat until we have filled up the **Feature Map**.

$$\begin{aligned} & (0 \times 0) + (0 \times 0) + (0 \times 1) \\ & + (1 \times 0) + (0 \times 1) + (0 \times 0) \\ & + (0 \times 1) + (1 \times 0) + (1 \times 0) \end{aligned}$$

$$= 0$$

+ -2

Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	-2



Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1
0	1	0
1	0	0

Then we shift the **Filter** over again and repeat until we have filled up the **Feature Map**.

$$\begin{aligned} & (0 \times 0) + (0 \times 0) + (0 \times 1) \\ & + (0 \times 0) + (0 \times 1) + (1 \times 0) \\ & + (1 \times 1) + (1 \times 0) + (0 \times 0) \end{aligned}$$

$$= 0$$

$$+ -2$$

Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	-1



Input Image

0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1
0	1	0
1	0	0

Then we shift the **Filter** over again and repeat until we have filled up the **Feature Map**.

$$\begin{aligned} & (0 \times 0) + (0 \times 0) + (1 \times 1) \\ & + (0 \times 0) + (1 \times 1) + (0 \times 0) \\ & + (1 \times 1) + (0 \times 0) + (0 \times 0) \end{aligned}$$

$$= 3$$

$$+ -2$$

Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



Input Image

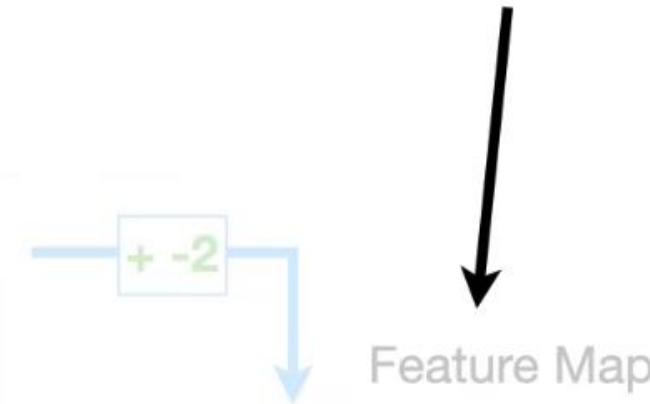
0	0	1	1	0	0
0	1	0	0	1	0
1	0	0	0	0	1
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0

Filter (aka Kernel)

0	0	1
0	1	0
1	0	0



BAM!!! We've filled in the whole **Feature Map**.



Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



Feature Map

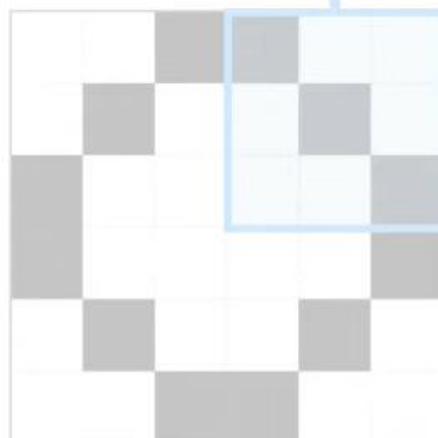
1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1

Feature Map, Post ReLU

1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1

Filter

(Convolution)



So let's move this stuff over to make some room for the next step.



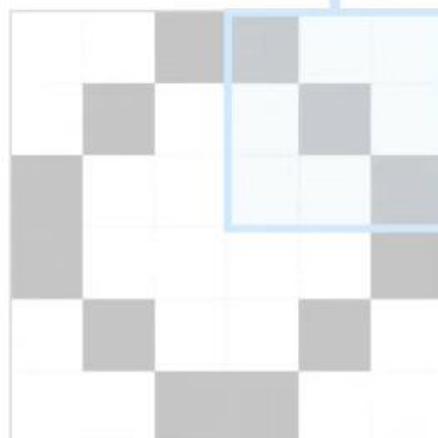
Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1

Feature Map, Post ReLU

1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1

Filter
(Convolution)



Now we apply another filter
to the new **Feature Map**.



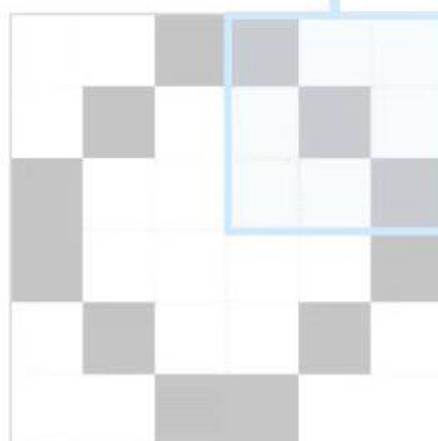
Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1

Feature Map, Post ReLU

1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1

Filter
(Convolution)



However, unlike before, we simply select the maximum value...



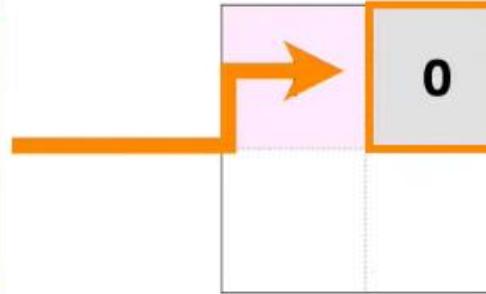
Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1

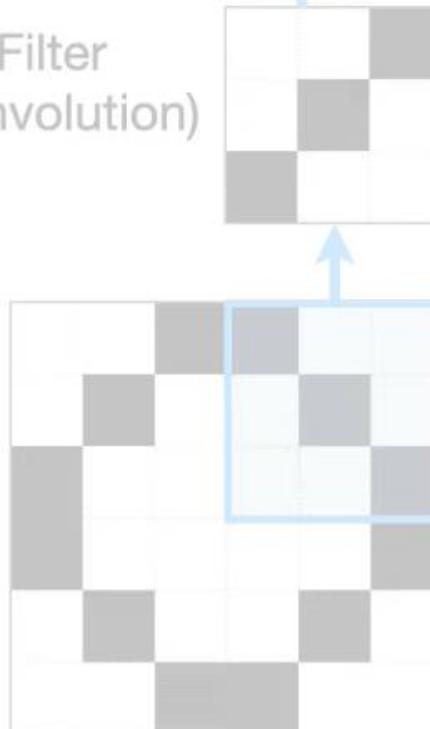


Feature Map, Post ReLU

1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1



Filter
(Convolution)



...and this filter usually moves in such a way that it does not overlap itself.



Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1

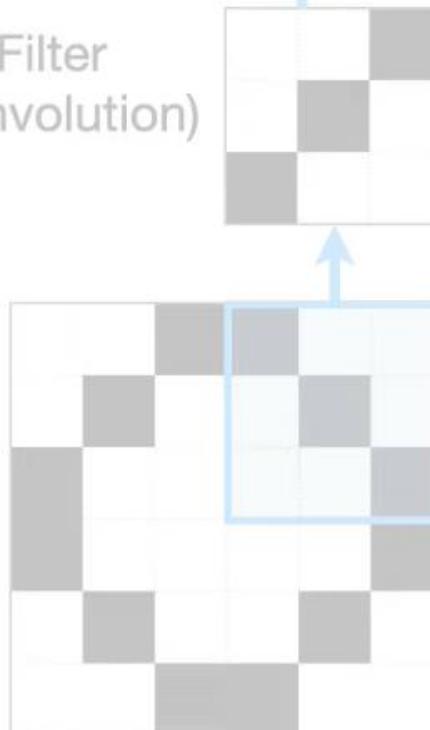


Feature Map, Post ReLU

1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1

1	0
0	

Filter
(Convolution)



...and this filter usually moves in such a way that it does not overlap itself.



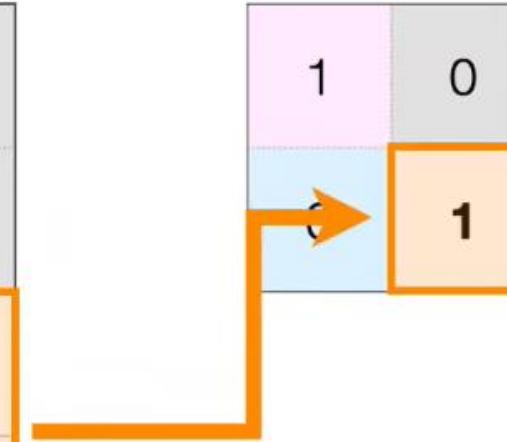
Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1

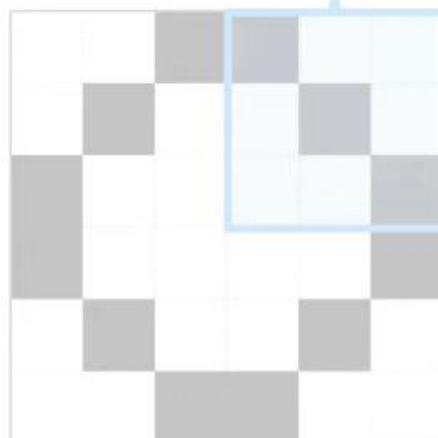


Feature Map, Post ReLU

1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1



Filter
(Convolution)



...and this filter usually moves in such a way that it does not overlap itself.



Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1

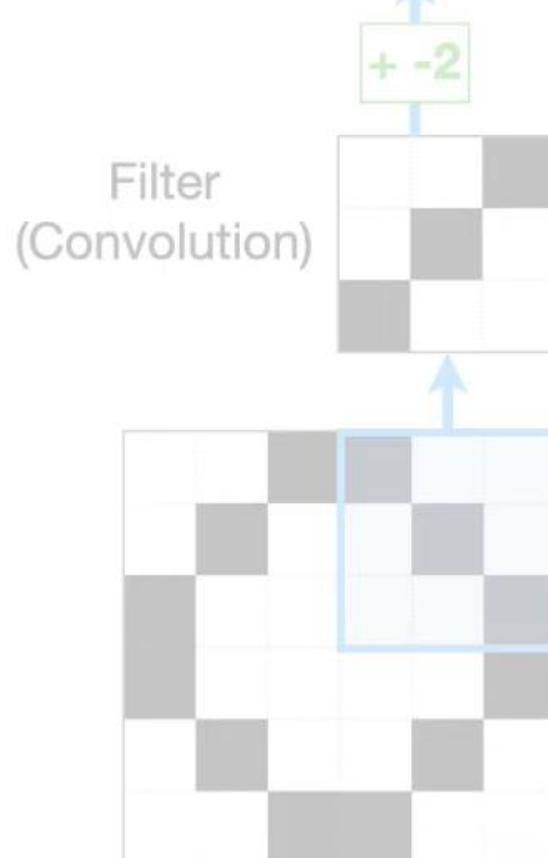


Feature Map, Post ReLU

1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1

Max Pooled

1	0
0	1



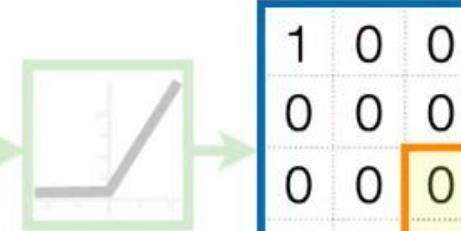
Terminology Alert!!!

When we select the maximum value in each region, we are applying **Max Pooling**.



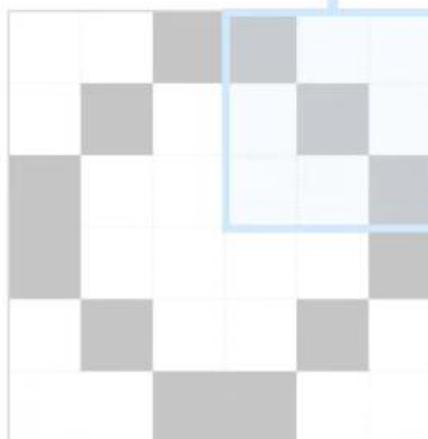
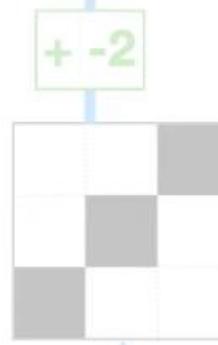
Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1

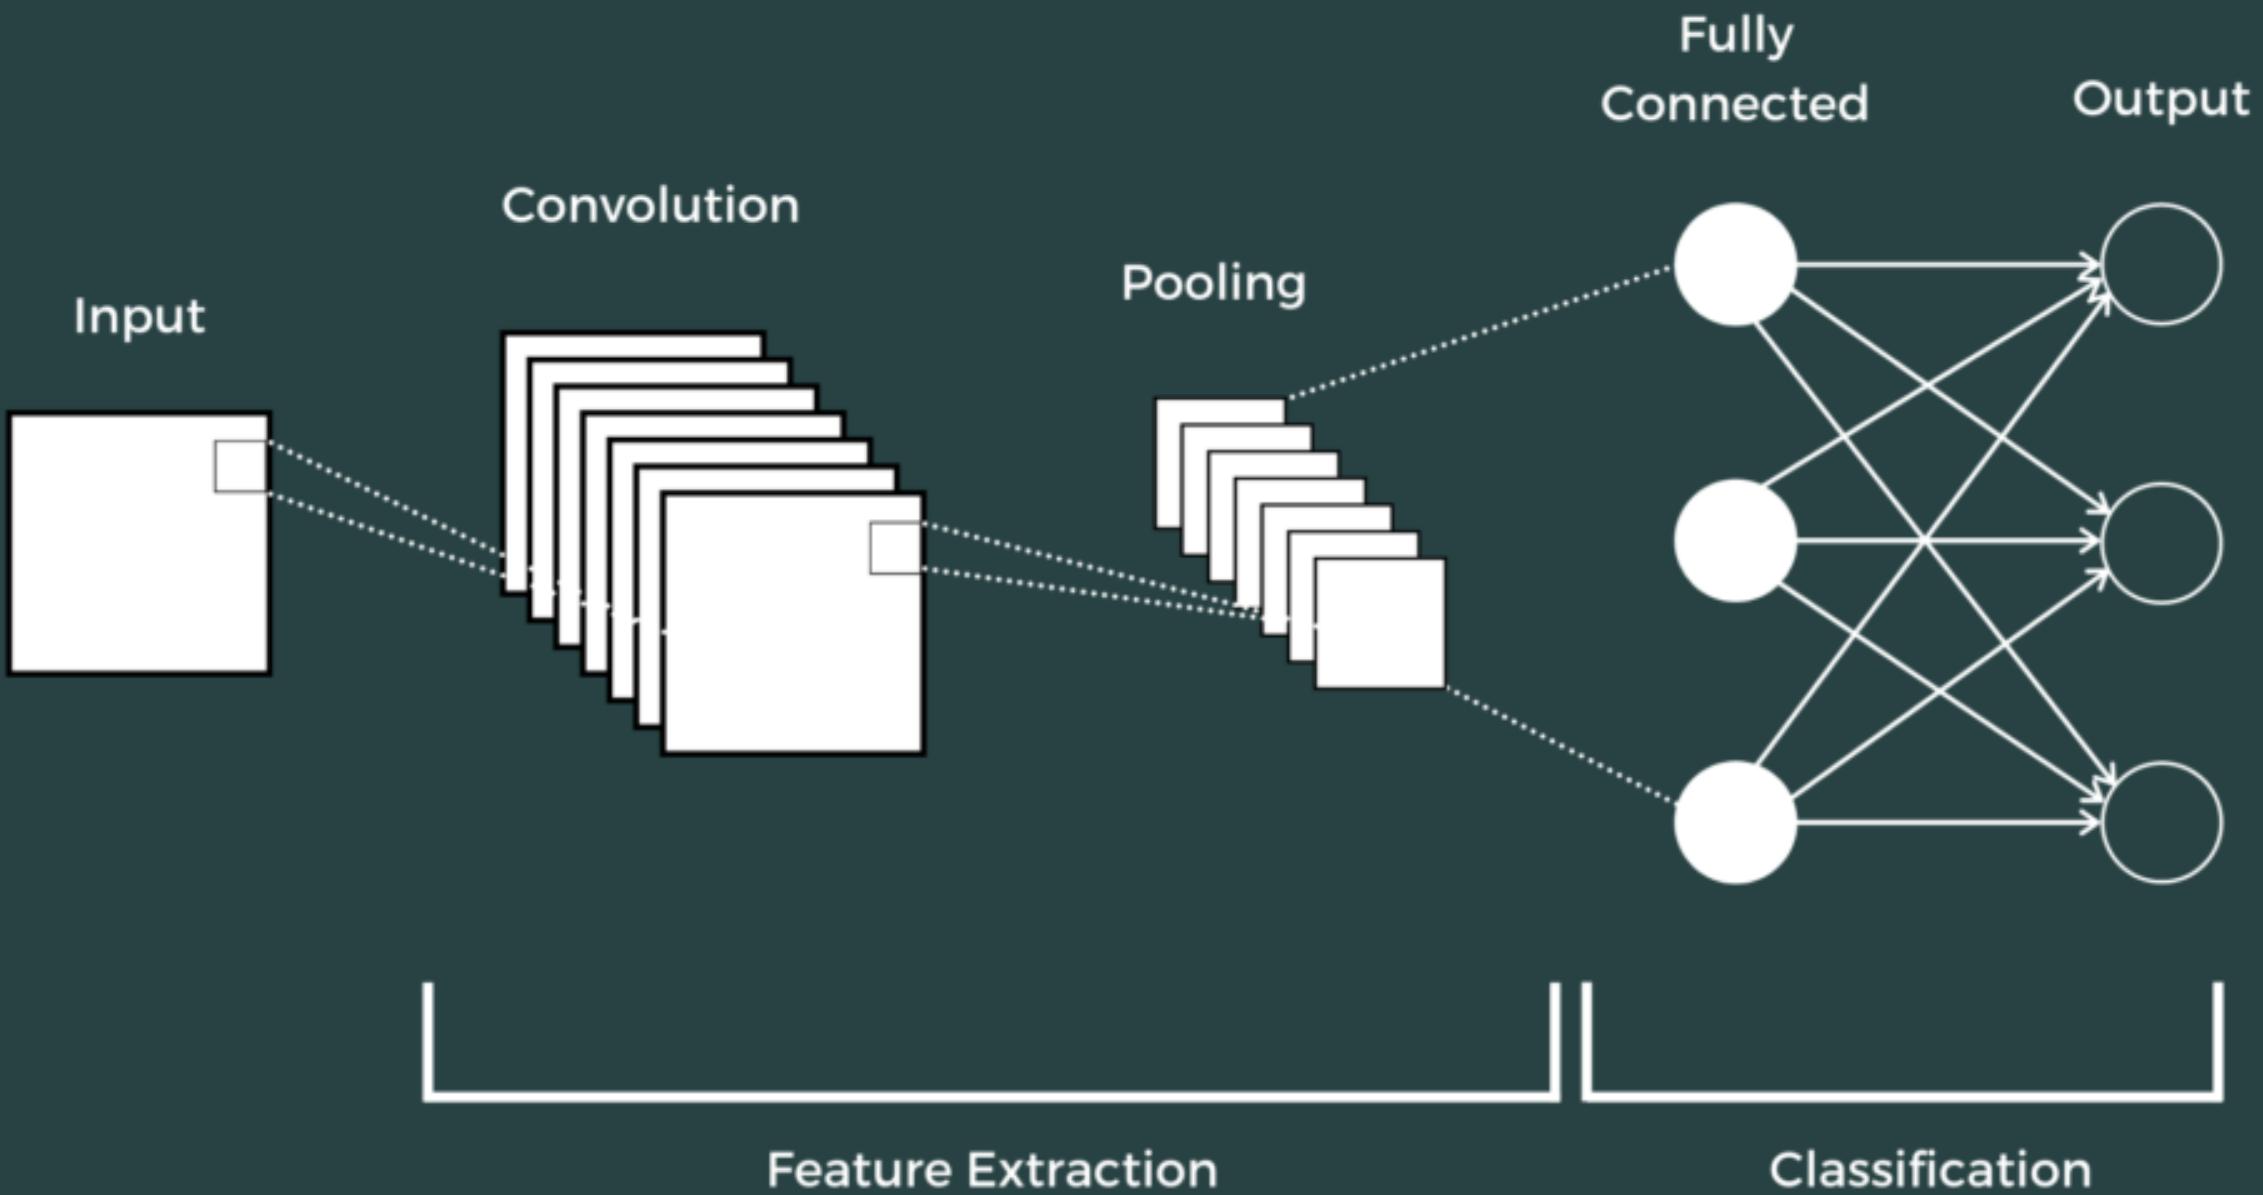


1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1

Filter
(Convolution)



...let's move and shrink things
to give us more room.





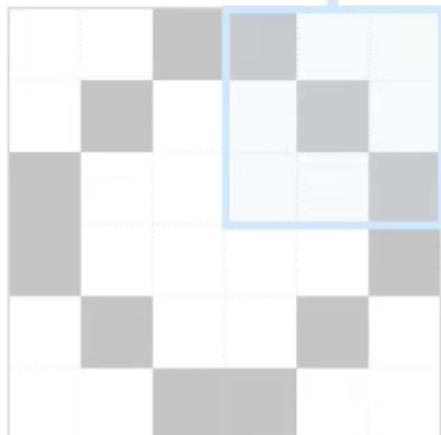
Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1

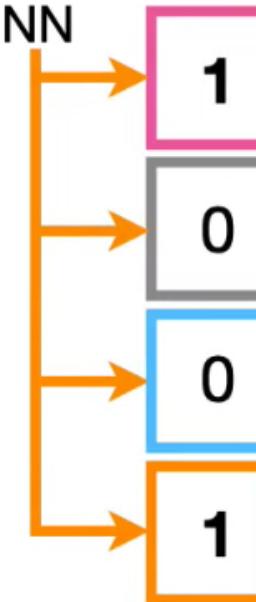


1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1

Filter
(Convolution)



Input to
NN



Now let's convert the
Pooled Layer into a
column of **Input Nodes**.



Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1

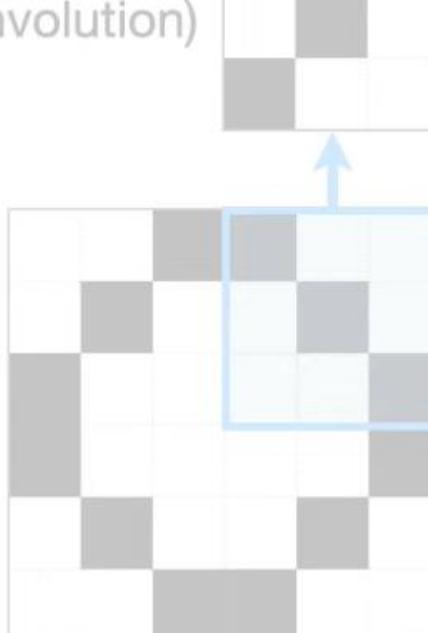


1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1

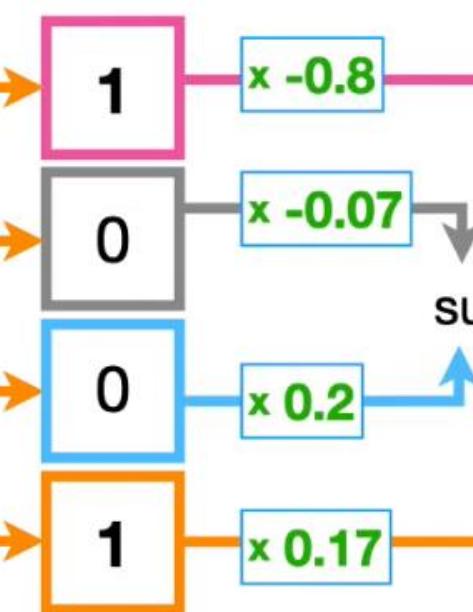
Max Pooling

1	0
0	1

Filter
(Convolution)

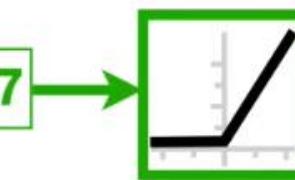


Input to
NN



sum

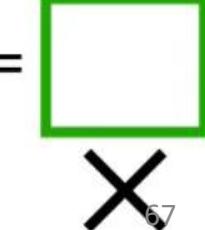
+ 0.97



+ 1.45 =



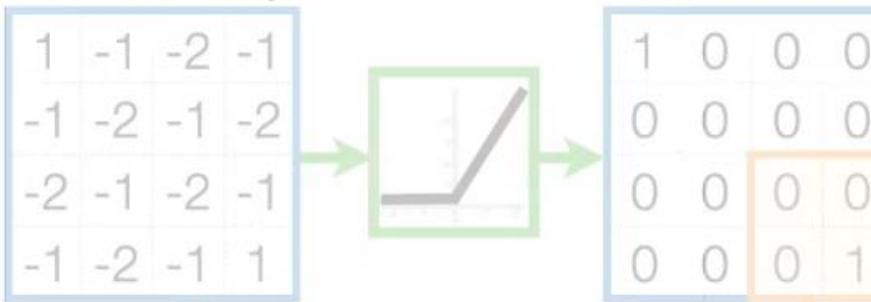
+ -0.45 =



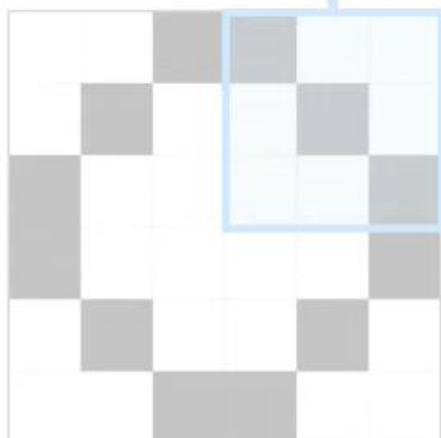
Lastly, let's plug the **Input
Nodes** into a normal, every day
Neural Network.



Feature Map

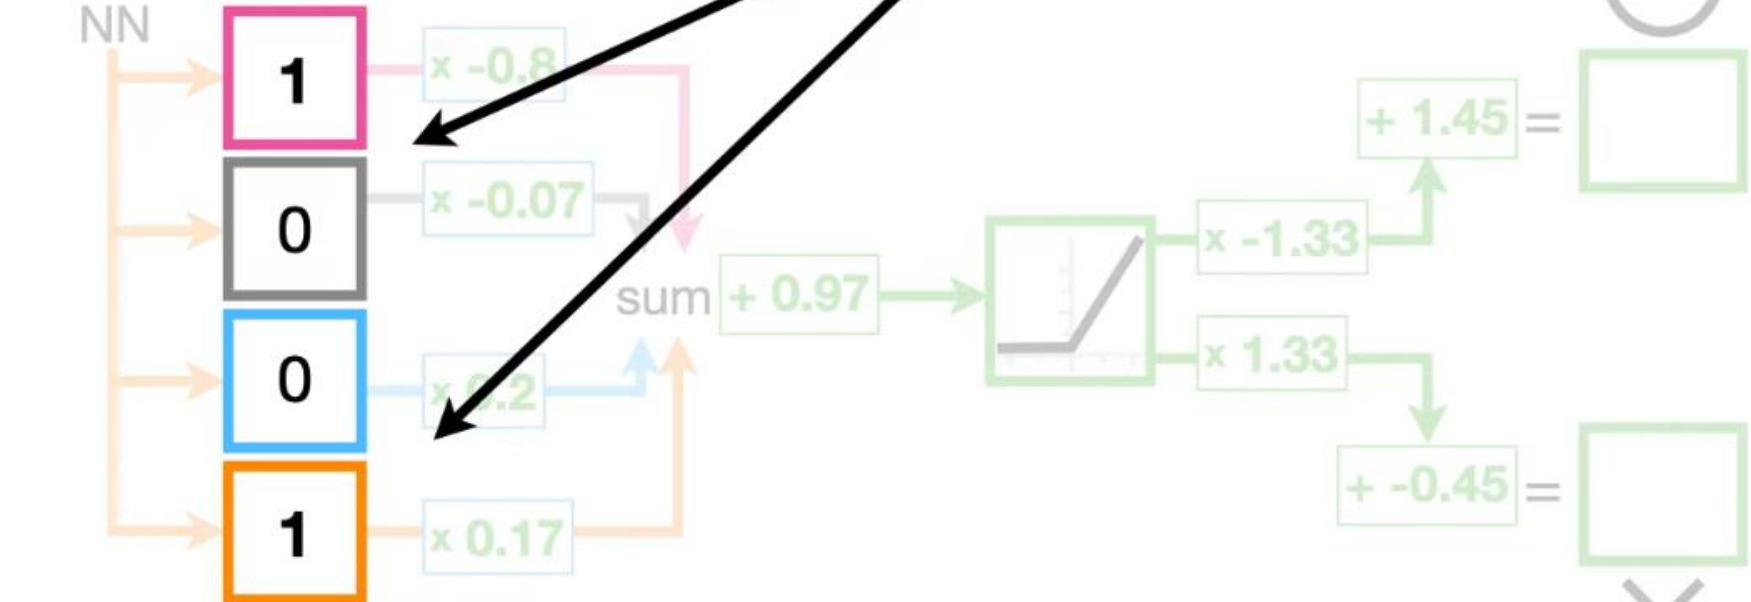


Filter
(Convolution)



Input to
NN

**This Neural Network has
4 Input Nodes...**





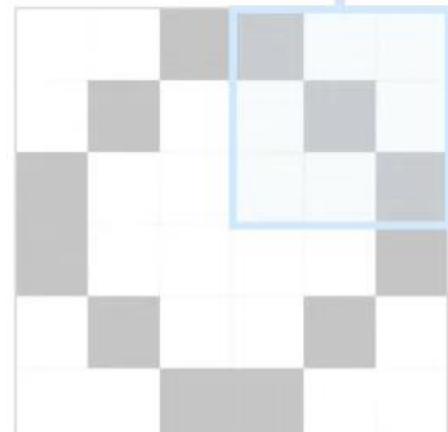
Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1

Filter
(Convolution)



1	0
0	1

Input to
NN

1

0

0

1

$\times -0.8$

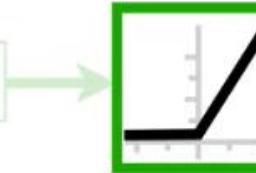
$\times -0.07$

$\times 0.2$

$\times 0.17$

sum

+ 0.97



$$+ 1.45 =$$

$$\times -1.33$$

$$\times 1.33$$

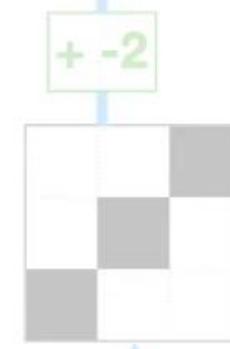
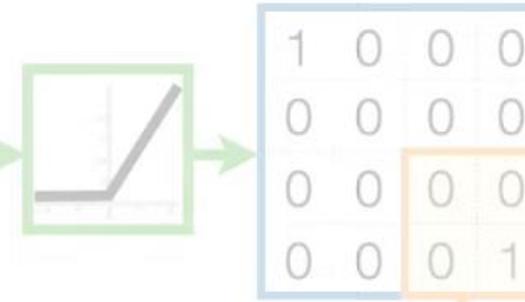
$$+ -0.45 =$$

...a single **Hidden Layer** with
a single **Node** using the
ReLU Activation Function...

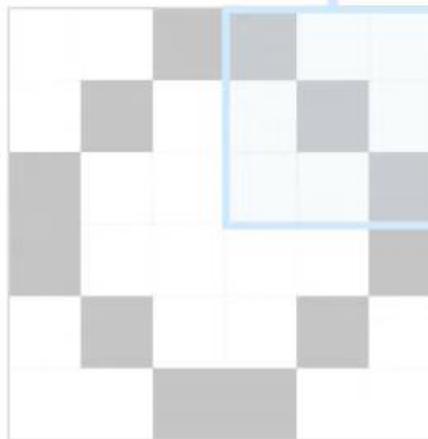


Feature Map

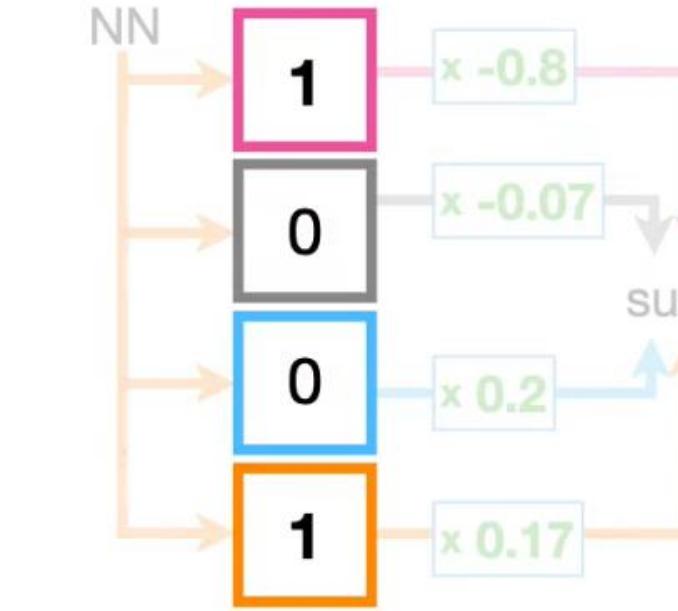
1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



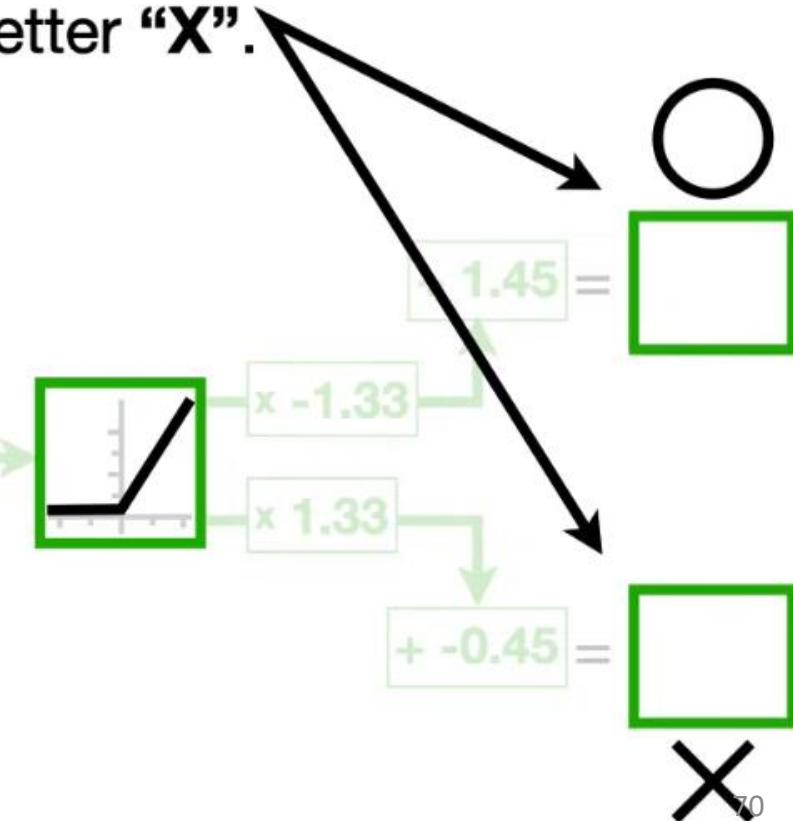
Filter
(Convolution)



Input to
NN



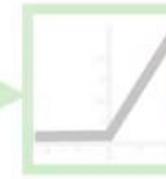
...and two **Output
Nodes**, one for the letter
“O” and one for the
letter “X”.





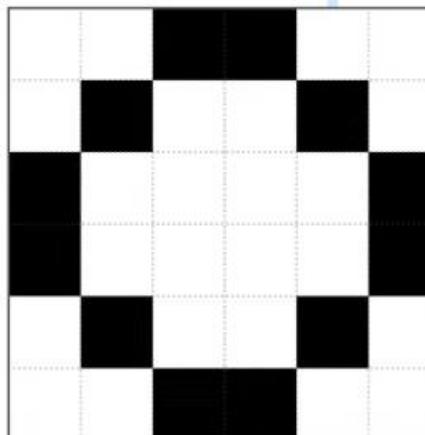
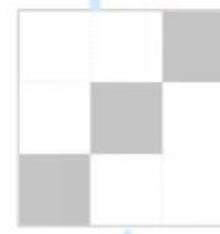
Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1

Filter
(Convolution)



1	0	
0	1	

Input to
NN

1

0

0

1

$\times -0.8$

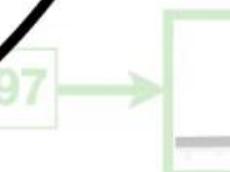
$\times -0.07$

$\times 0.2$

$\times 0.17$

sum

$+ 0.97$



So, given this image...





Feature Map

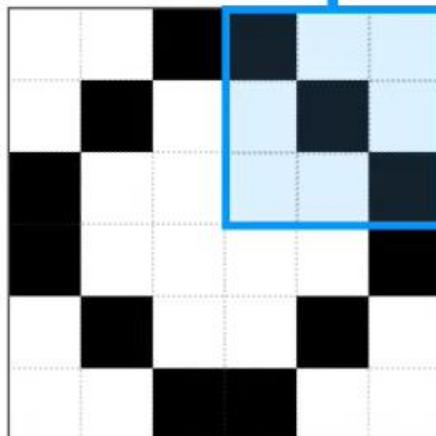
1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



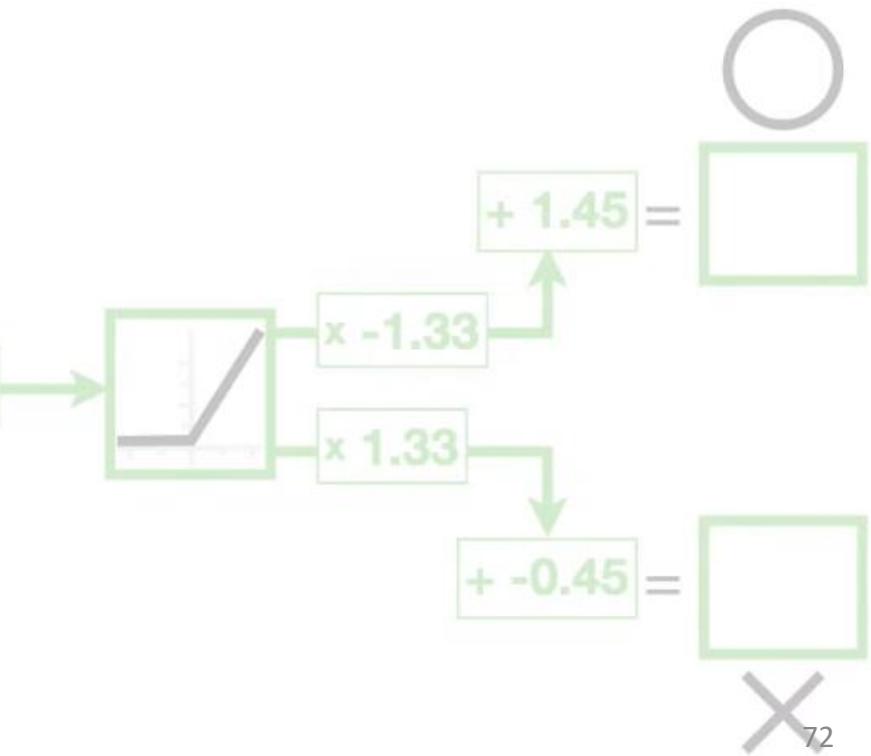
1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1

...we run it through
the **Filter** to create
the **Feature Map**...

Filter
(Convolution)



Input to
NN

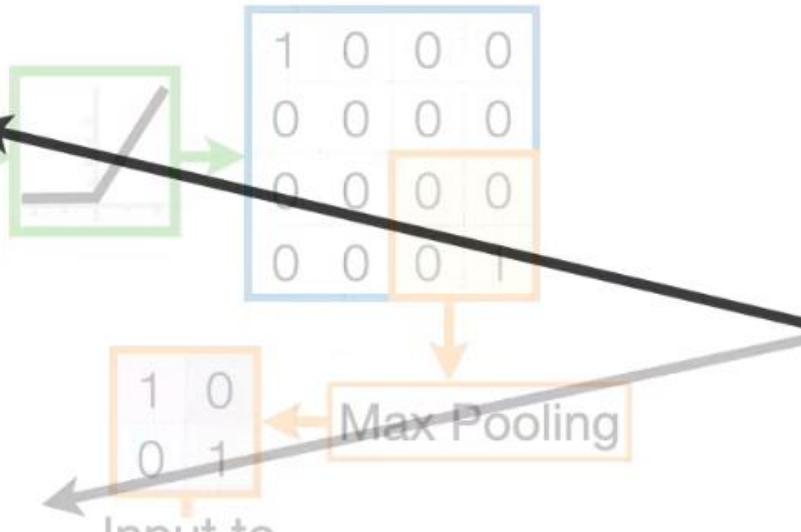
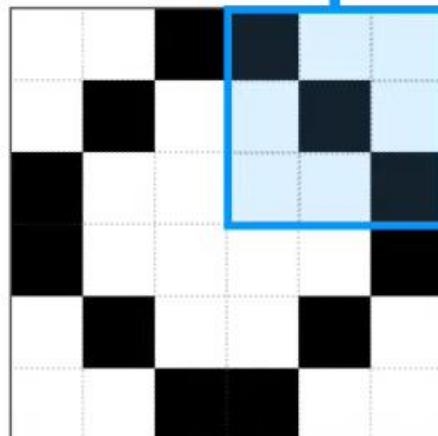




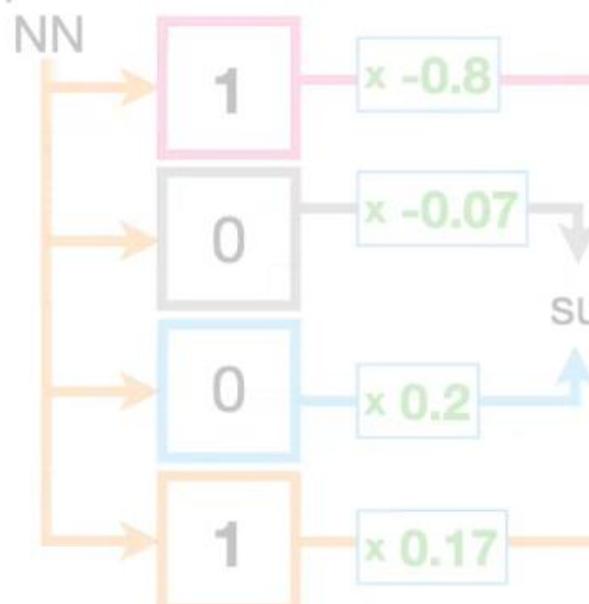
Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1

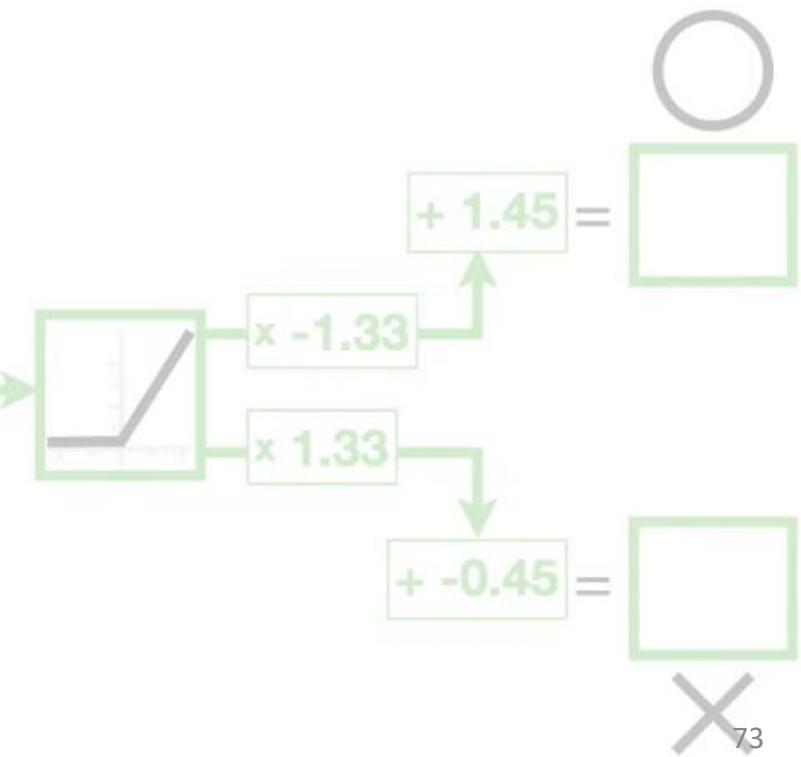
Filter
(Convolution)



Input to NN



...we run it through
the **Filter** to create
the **Feature Map**...





Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



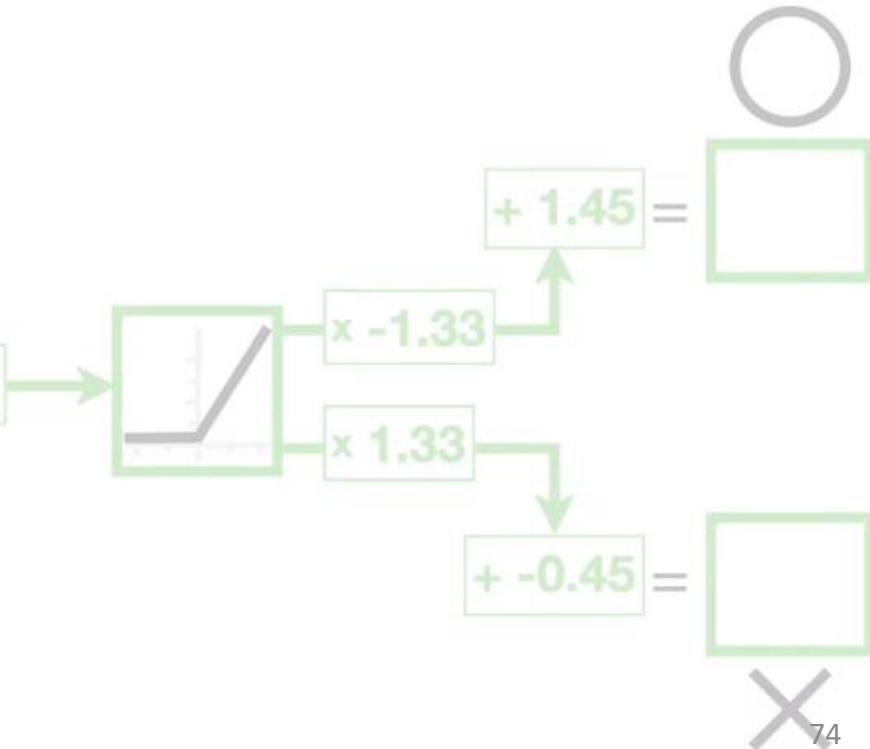
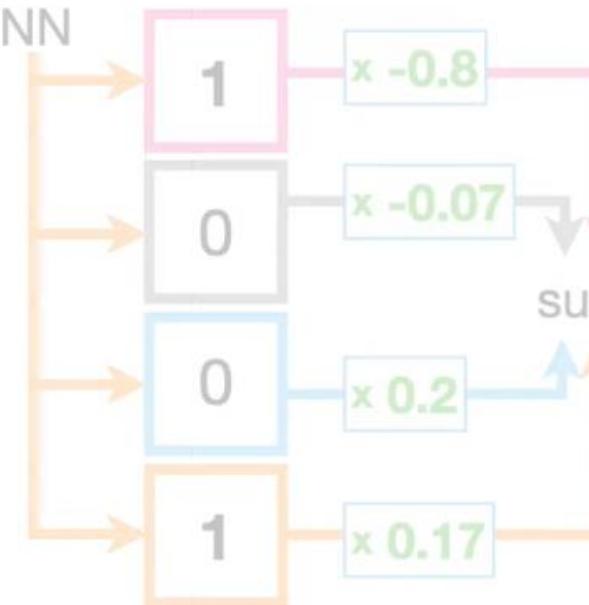
1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1

...then we run the
Feature Map through
**a ReLU Activation
Function...**

Filter
(Convolution)

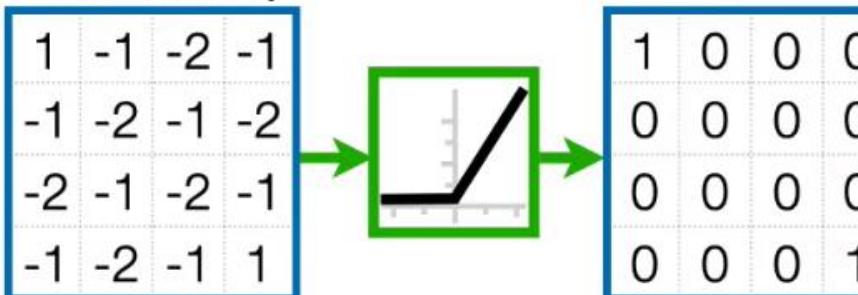
1	0
0	1

Input to
NN

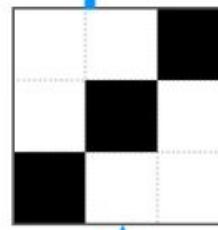




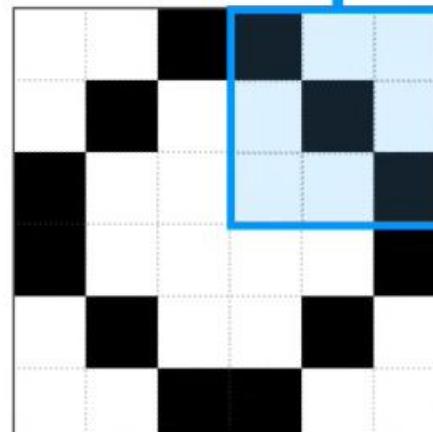
Feature Map



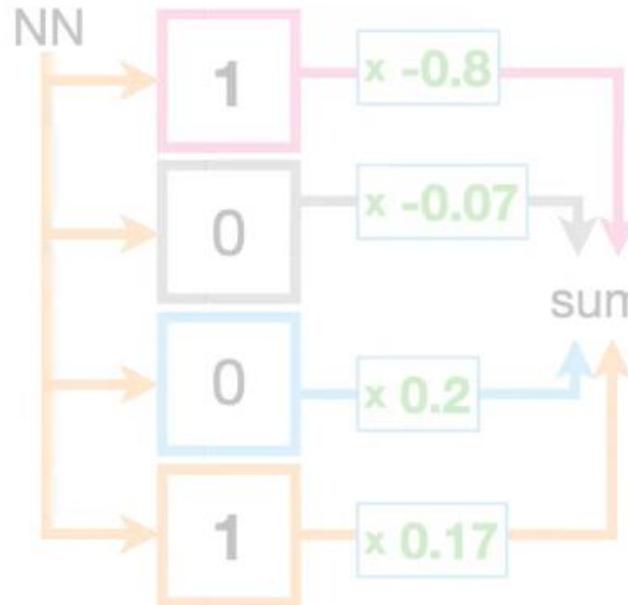
+ -2



Filter
(Convolution)



Input to
NN



...then we run the
Feature Map through
**a ReLU Activation
Function...**





Feature Map

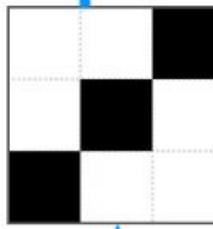
1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1

...then we select the maximum value in each area...

Filter
(Convolution)



+ -2

1	0
0	1

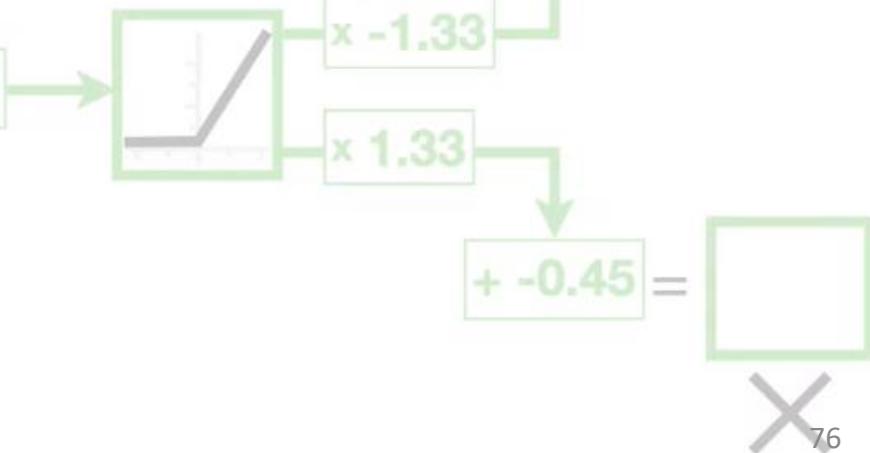
Max Pooling

Input to NN



sum

+ 0.97





Feature Map

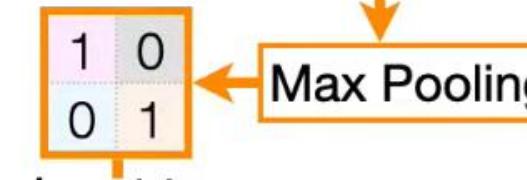
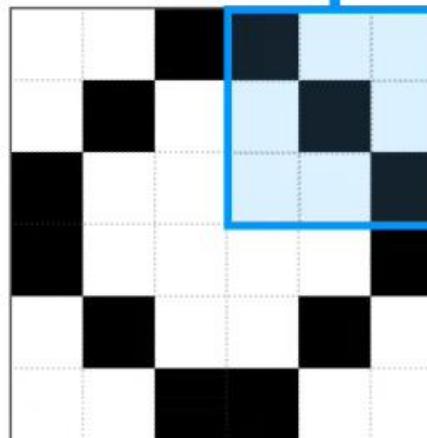
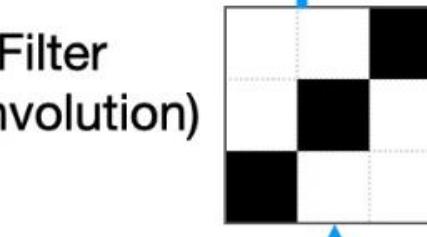
1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1

... and end up with these values in the **Input Nodes**.

Filter
(Convolution)



Input to
NN

1

0

0

1

$\times -0.9$

$\times -0.07$

$\times 0.2$

$\times 0.17$

\sum

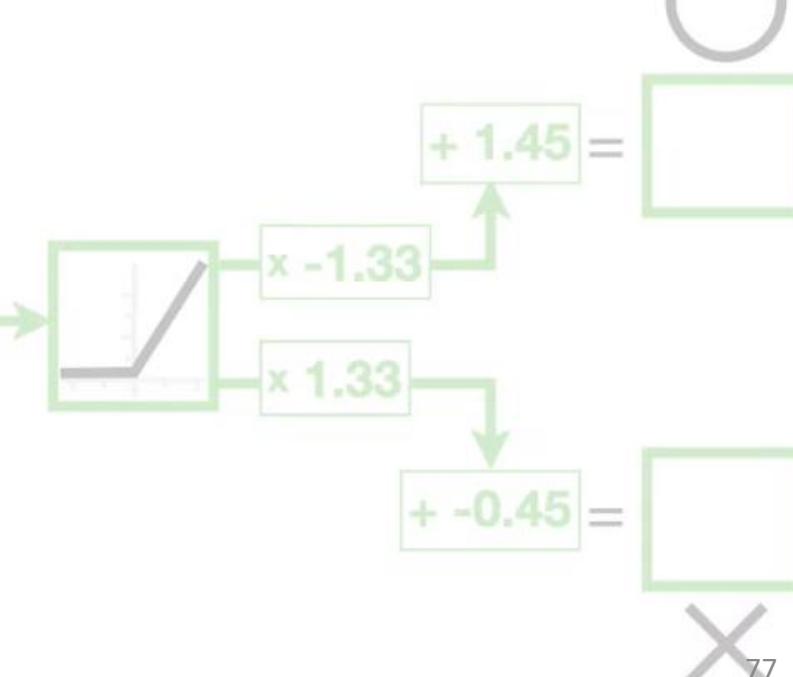
$+ 0.97$

\rightarrow

$\times -1.33$

$\times 1.33$

$+ -0.45$



Double BAM!!
SQ!

Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



$$(1 \times -0.8) \quad (0 \times -0.07)$$

0	0	0
0	0	0
0	0	0.1

1	0
0	1

Max Pooling

Now we multiply the values in the **Input Nodes** by their associated **Weights**...

Filter
(Convolution)



Input to NN

1	x -0.8
0	x -0.07
0	x 0.2
1	x 0.17

$$x -0.8$$

$$x -0.07$$

$$x 0.2$$

$$x 0.17$$

sum

$$+ 0.97$$

$$x -1.33$$

$$x 1.33$$

$$+ 1.45 =$$



$$+ 1.45$$

$$x -1.33$$

$$x 1.33$$

$$+ -0.45$$



Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



$$(1 \times -0.8) + (0 \times -0.07)$$

0	0	0	0
0	0	0	0
0	0	0	1

$$(0 \times 0.2)$$

Now we multiply the values in the **Input Nodes** by their associated **Weights**...

Max Pooling

Input to NN

$$\begin{array}{c} 1 \\ \times -0.8 \\ \hline 0 \\ \times -0.07 \\ \hline 0 \\ \times 0.2 \\ \hline 1 \\ \times 0.17 \\ \hline \end{array}$$

sum

$$+ 0.97$$

$$\rightarrow$$

$$\times -1.33$$

$$\rightarrow$$

$$\times 1.33$$

$$\rightarrow$$

$$+ -0.45$$





Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1

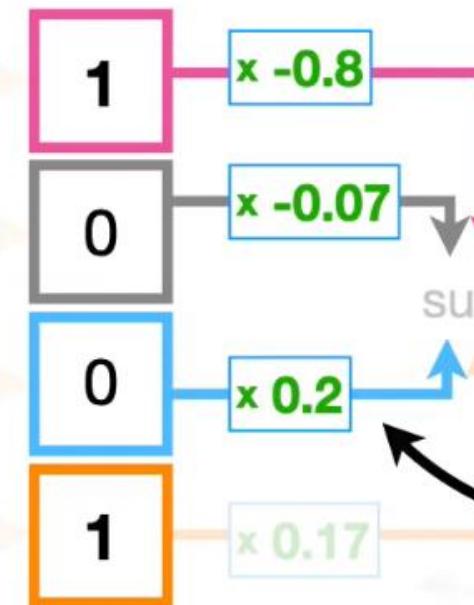


$$(1 \times -0.8) \quad (0 \times -0.07) \quad (0 \times 0.2)$$

0	0	0	0
0	0	0	0
0	0	0	1
0	0	0	1



Input to NN



Now we multiply the values in the **Input Nodes** by their associated **Weights**...

sum
+ 0.97

x -1.33

x 1.33

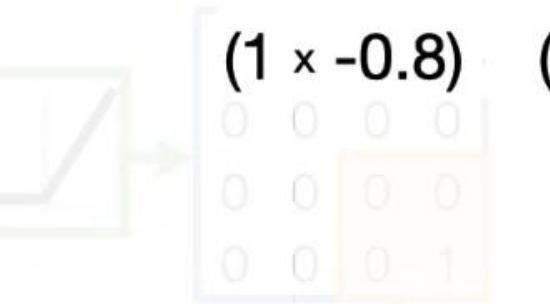
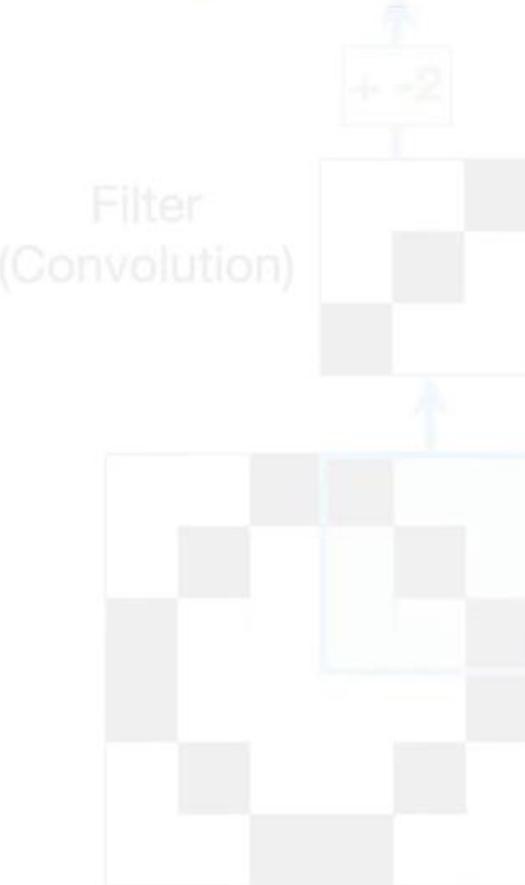
+ 1.45 =

+ -0.45 =

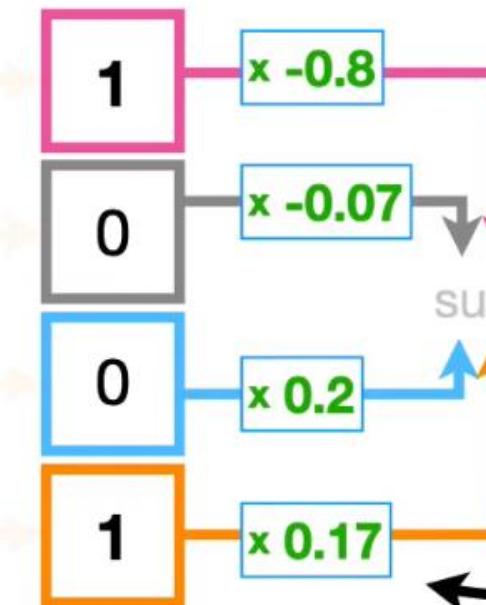


Feature Map

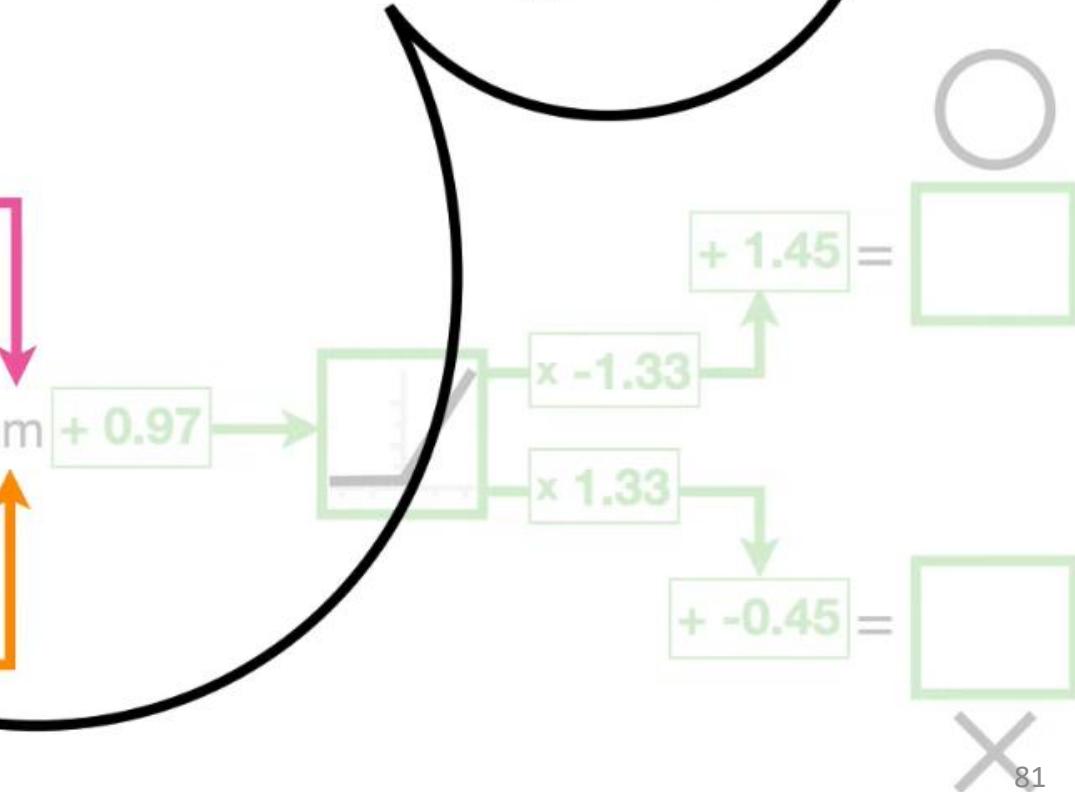
1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1

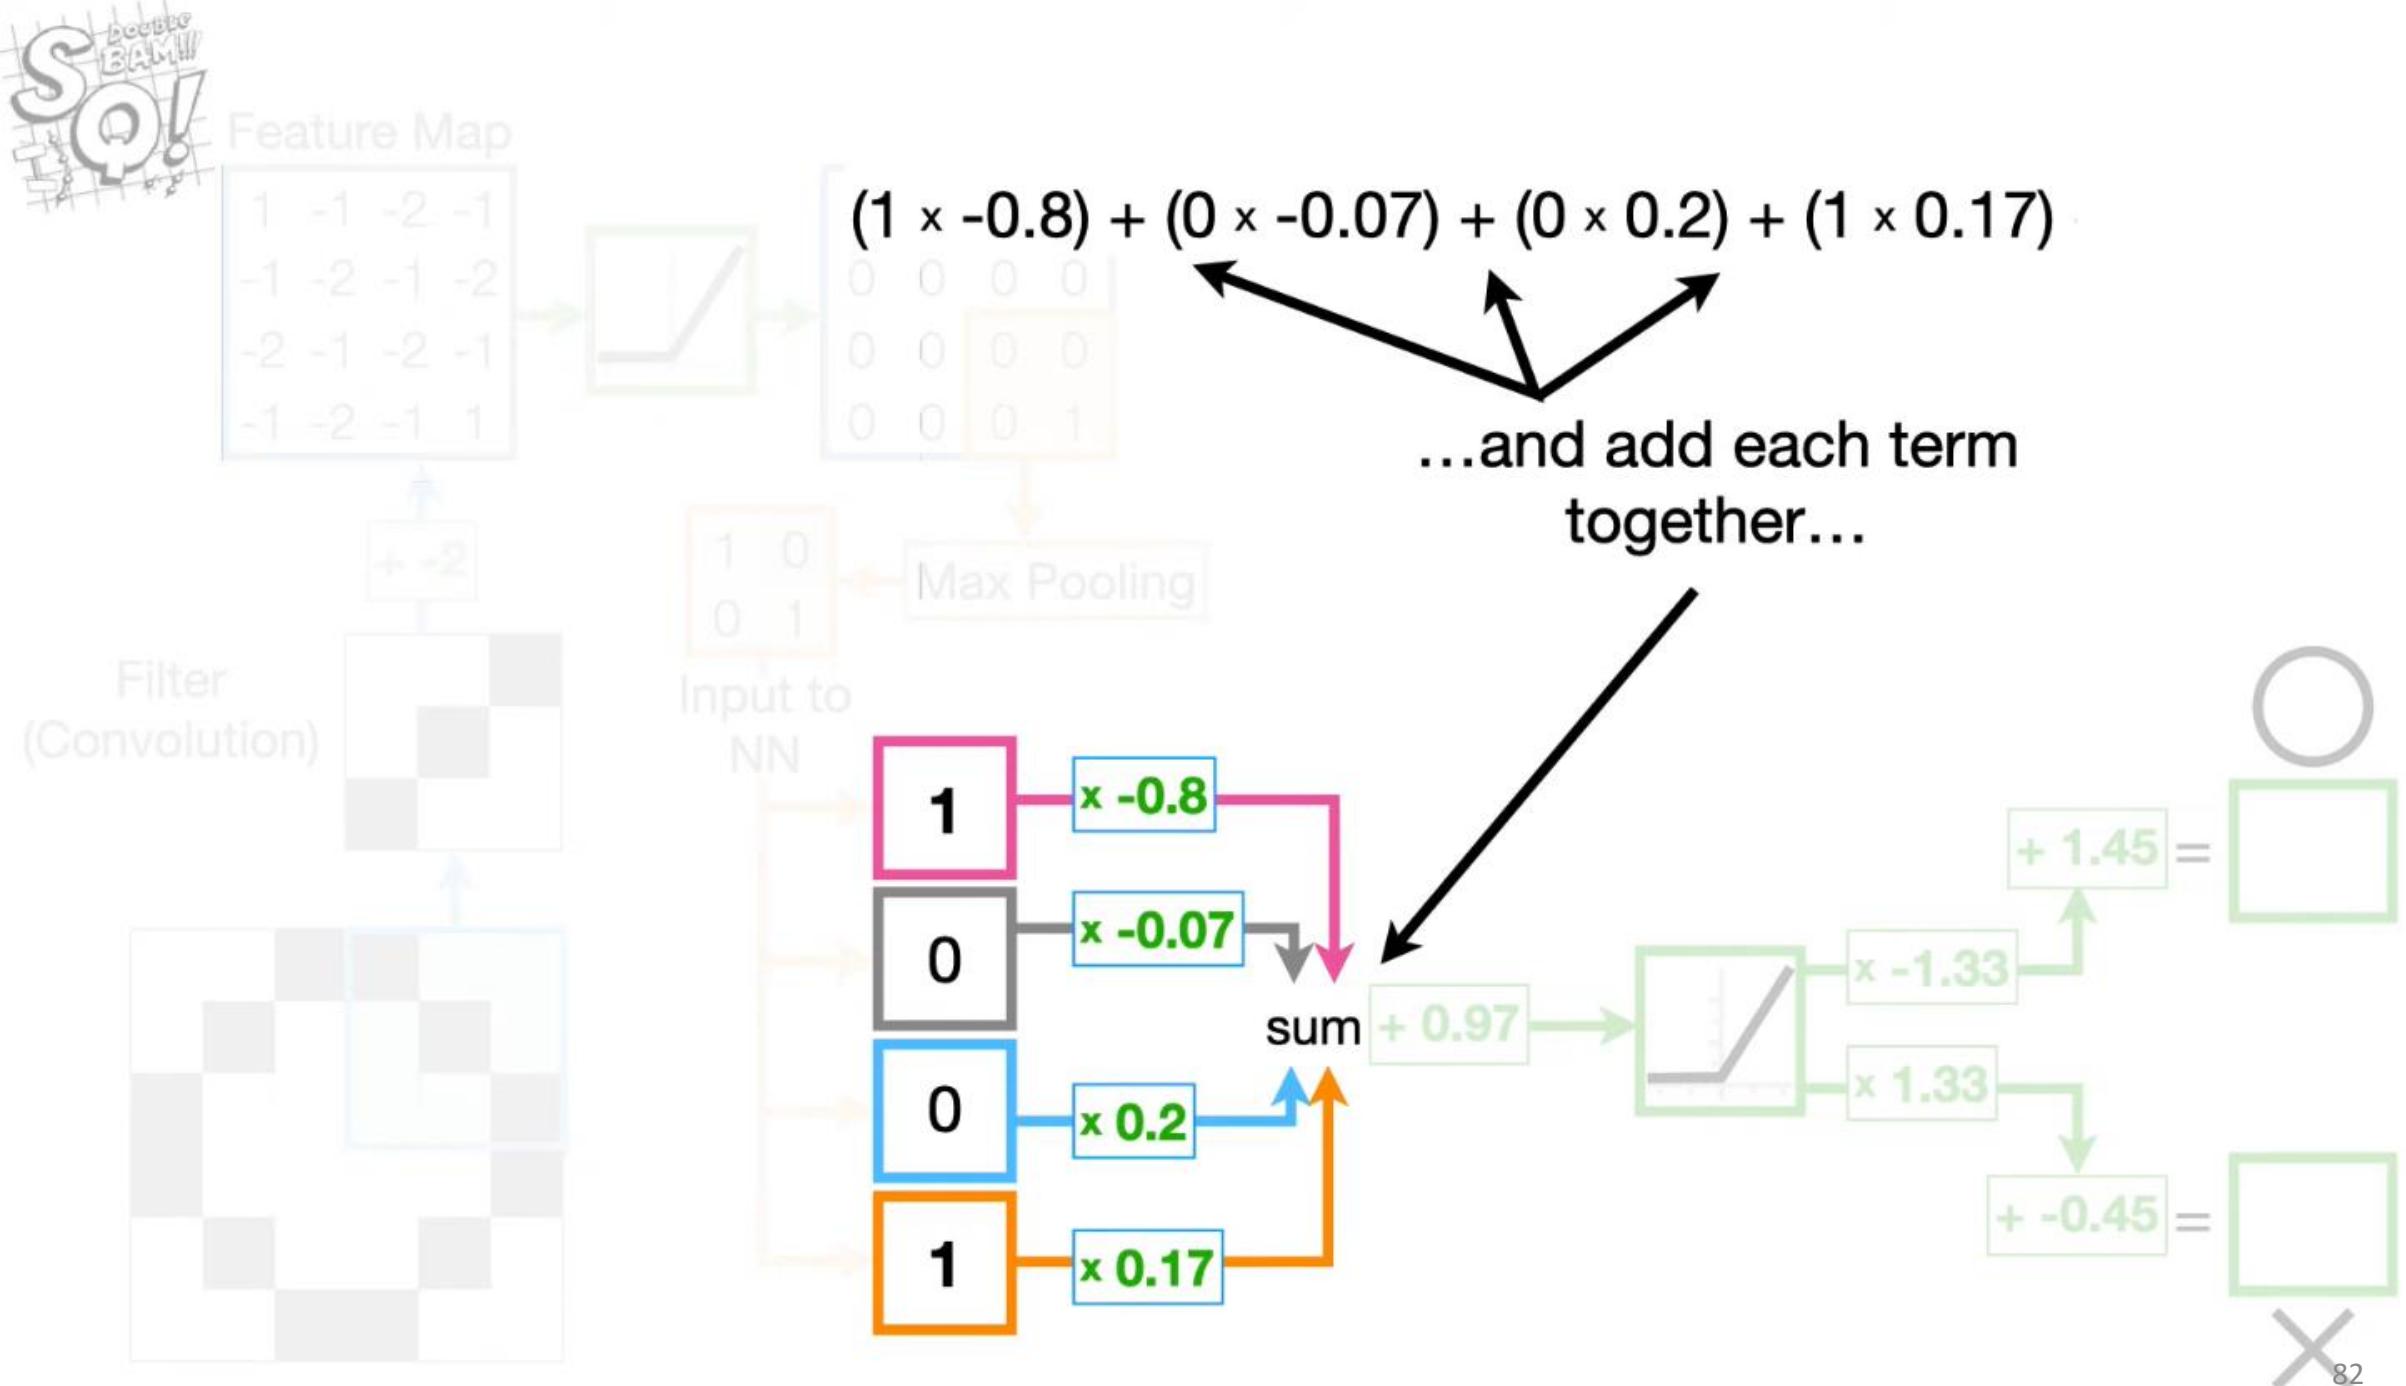


Input to NN



Now we multiply the values in the **Input Nodes** by their associated **Weights**...







Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



$$(1 \times -0.8) + (0 \times -0.07) + (0 \times 0.2) + (1 \times 0.17) + 0.97$$

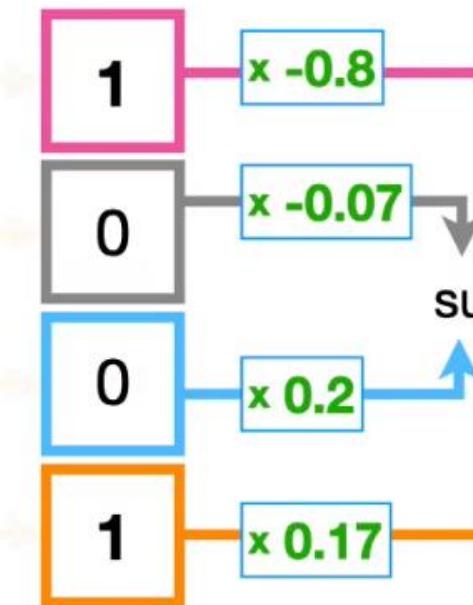
0	0	0	0
0	0	0	0
0	0	0	1



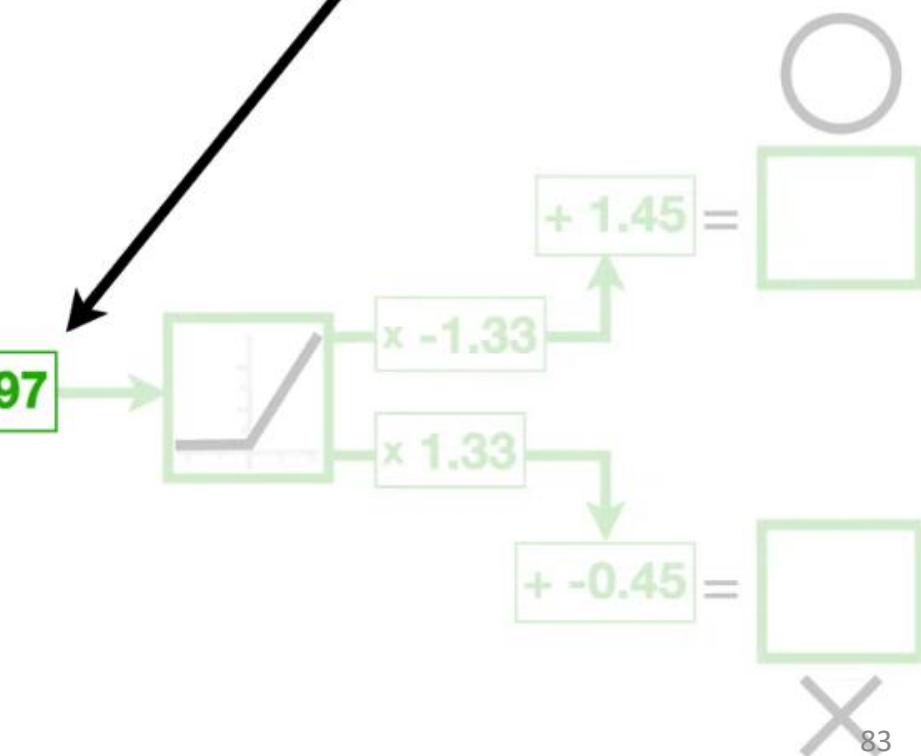
1 0
0 1

Max Pooling

Input to NN



...and then add the **bias**...





Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



$$(1 \times -0.8) + (0 \times -0.07) + (0 \times 0.2) + (1 \times 0.17) + 0.97 = 0.34$$

0	0	0	0
0	0	0	0
0	0	0	1

...and we end up with **0.34**.

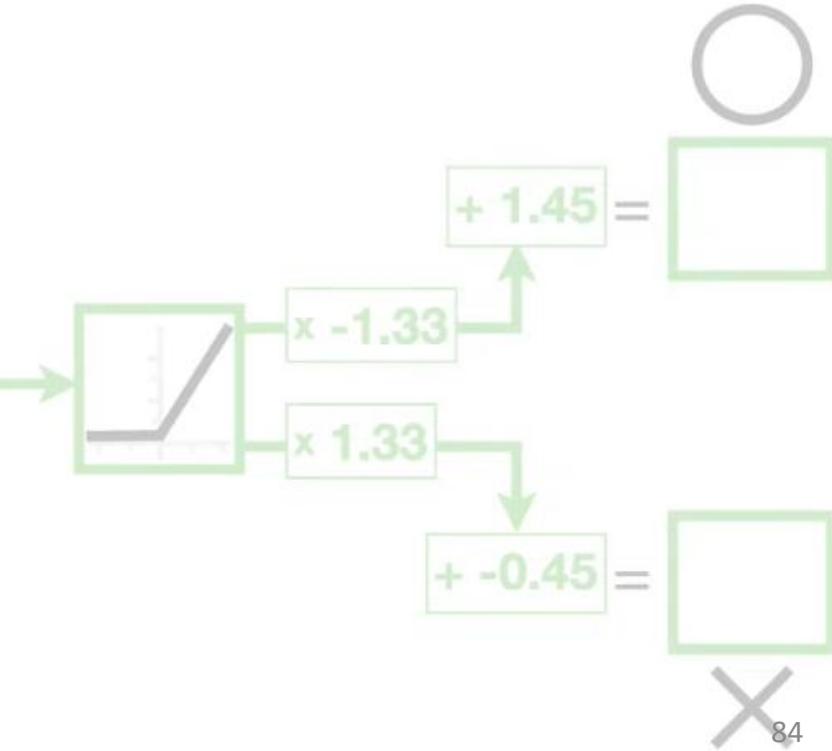
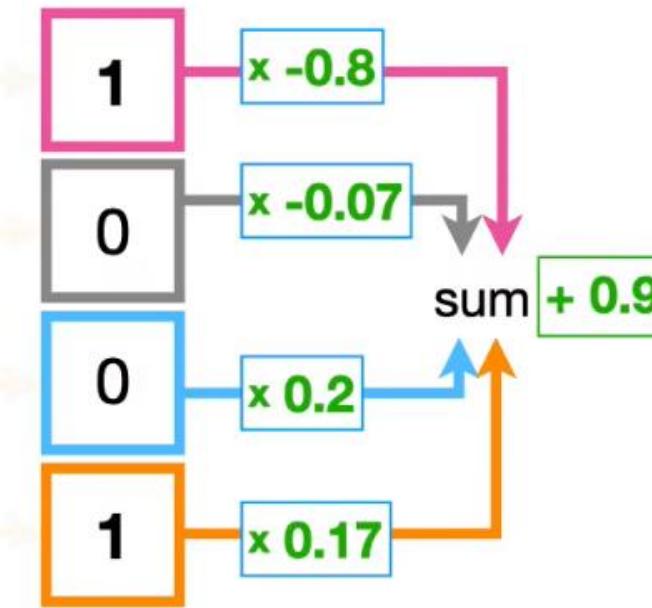
Filter
(Convolution)



1	0
0	1

Max Pooling

Input to
NN





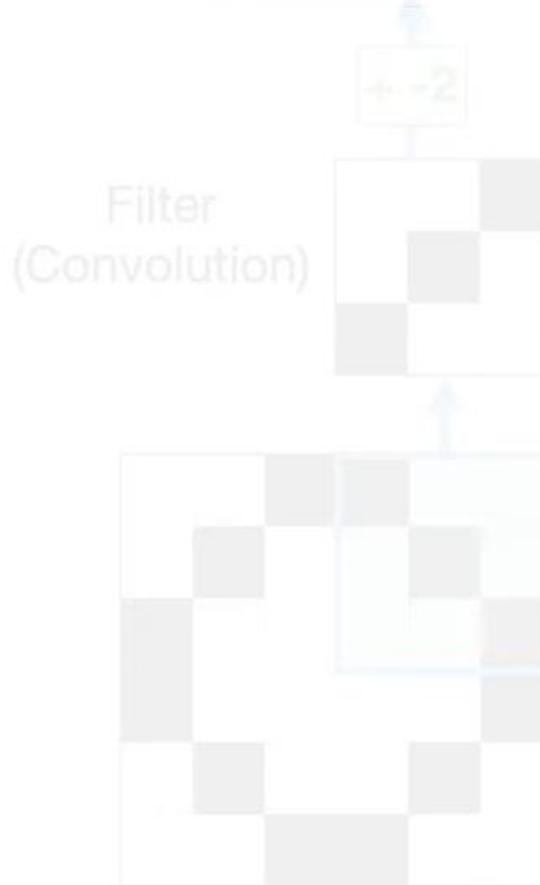
Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1

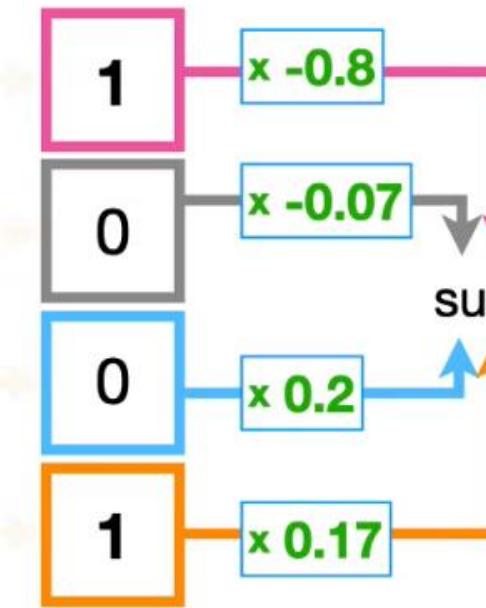


$$(1 \times -0.8) + (0 \times -0.07) + (0 \times 0.2) + (1 \times 0.17) + 0.97 = 0.34$$

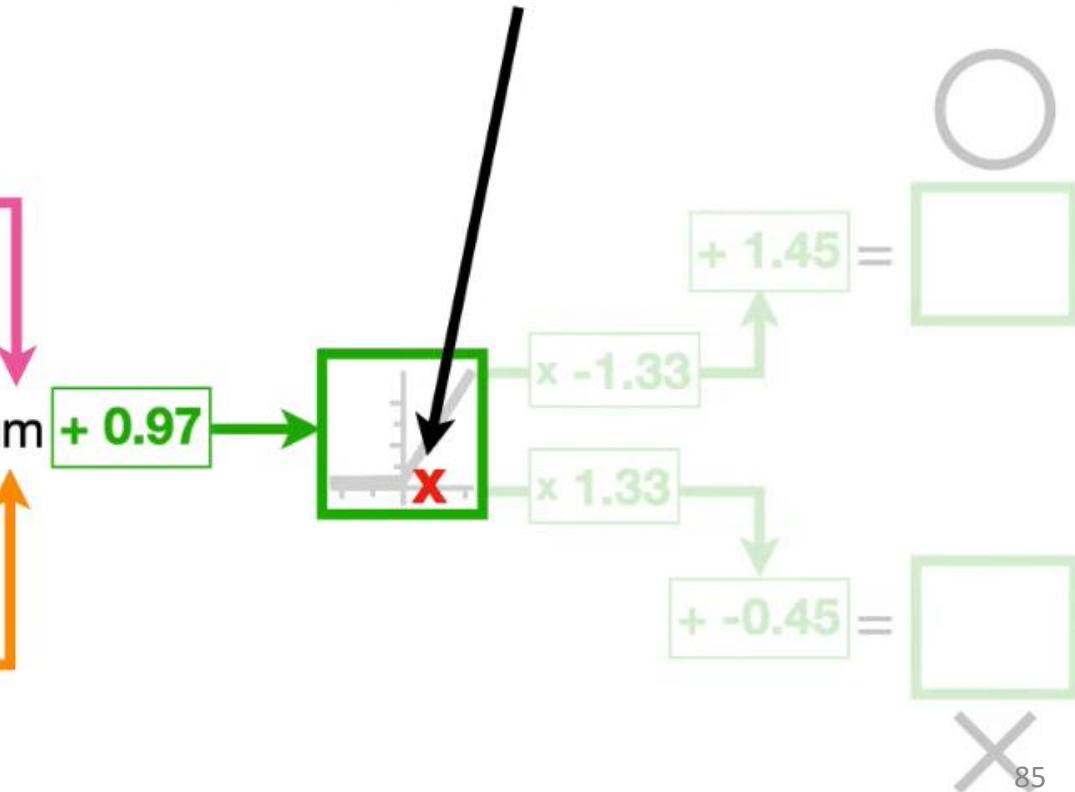
0	0	0	0
0	0	0	0
0	0	0	1



Input to NN



And thus, the x-axis coordinate for the **Activation Function** is **0.34**.





Feature Map

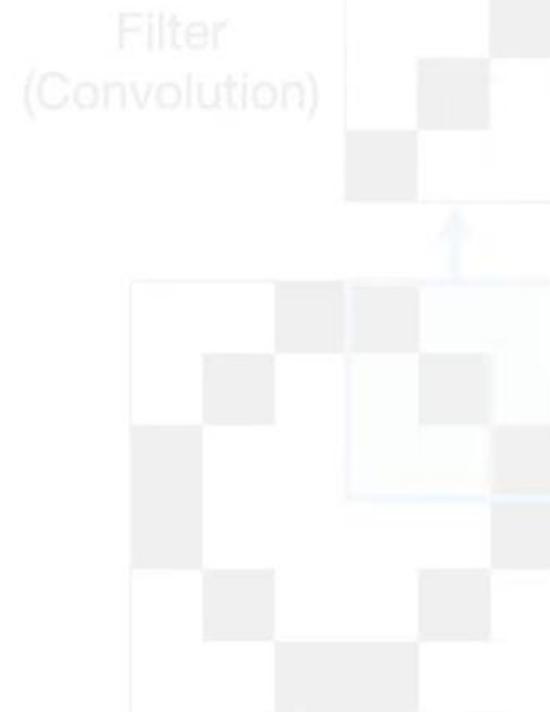
1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



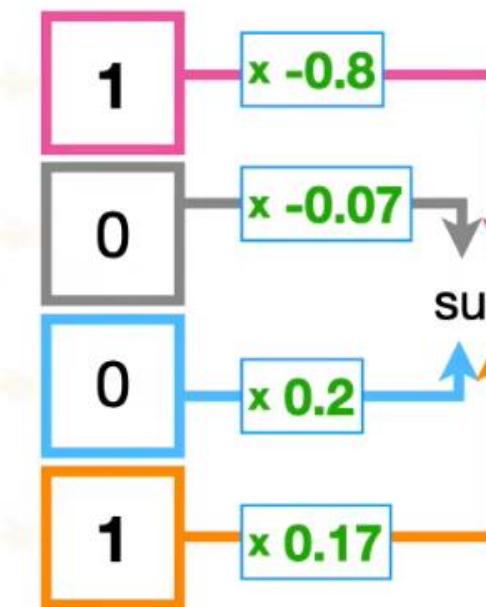
$$(1 \times -0.8) + (0 \times -0.07) + (0 \times 0.2) + (1 \times 0.17) + 0.97 = 0.34$$

Now we plug **0.34** into the **ReLU Activation Function**...

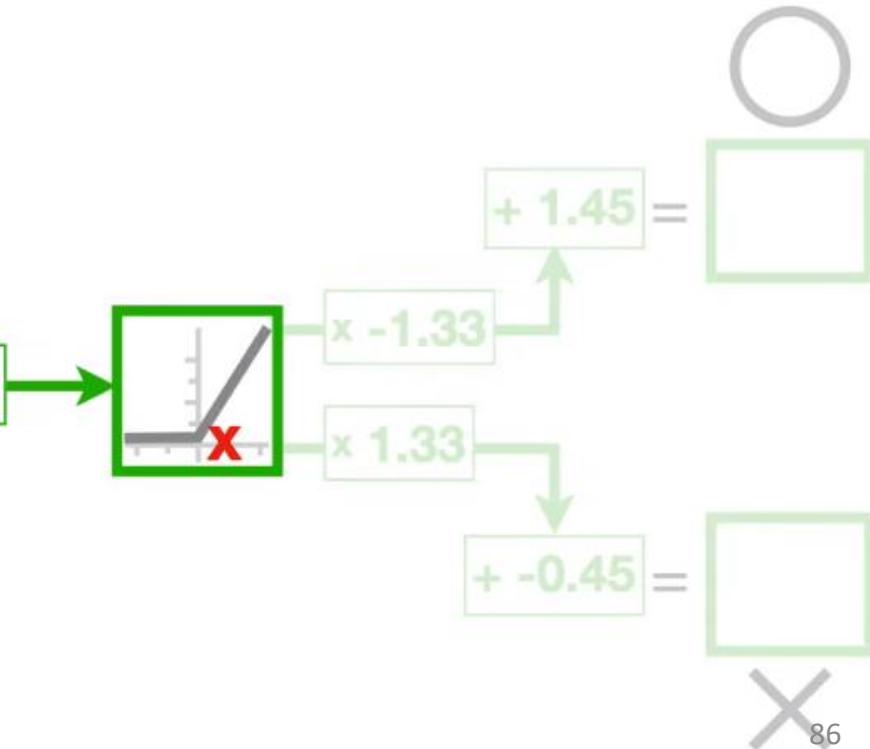
$$f(x) = \max(0, x) = y\text{-axis coordinate}$$



Input to NN



sum + 0.97





Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1

+ -2



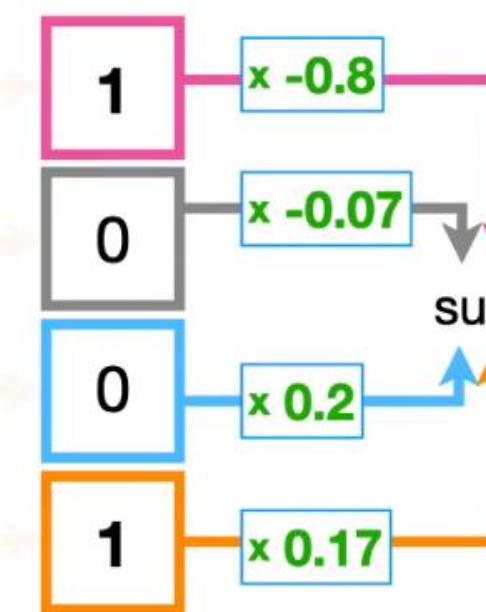
Filter
(Convolution)



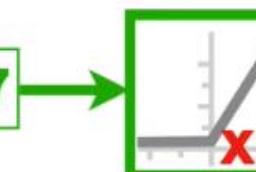
$$(1 \times -0.8) + (0 \times -0.07) + (0 \times 0.2) + (1 \times 0.17) + 0.97 = 0.34$$

Now we plug **0.34** into $f(0.34) = \max(0, 0.34) = y\text{-axis coordinate}$
the **ReLU Activation Function**...

Input to NN



sum + 0.97



+ 1.45 =



x -1.33

x 1.33

+ -0.45 =





Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1

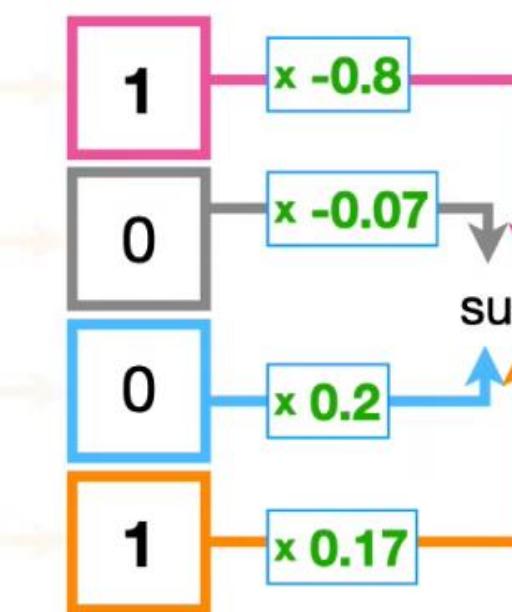


...and the output
is **0.34**, because

0.34 > 0.

Input to NN

Filter
(Convolution)



$$(1 \times -0.8) + (0 \times -0.07) + (0 \times 0.2) + (1 \times 0.17) + 0.97 = 0.34$$

$$f(0.34) = \max(0, 0.34) = 0.34$$





Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1

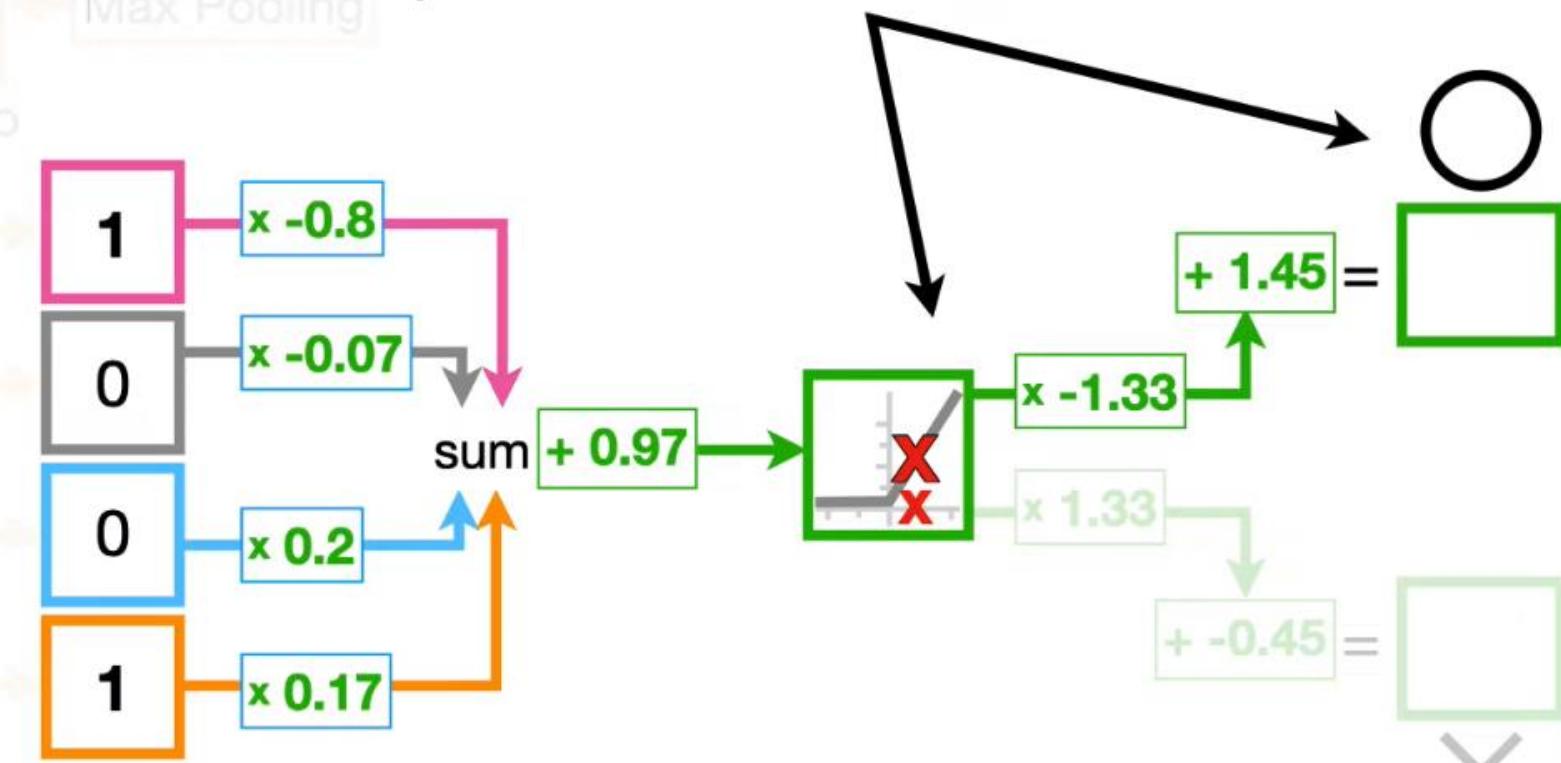
1	0
0	1

Max Pooling

Input to
NN

Filter
(Convolution)

Now, this connection, from
the **Hidden Layer** to the
output for the letter “O”...





Feature Map

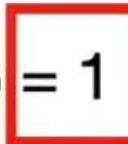
1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1

...gives us 1 for the letter “O”.

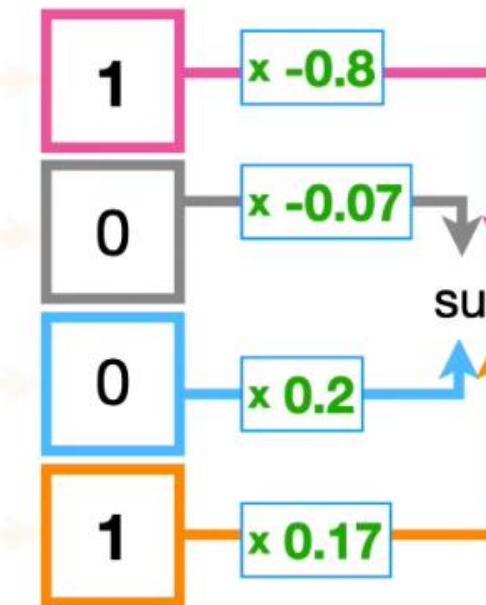
$$(0.34 \times -1.33) + 1.45 = 1$$



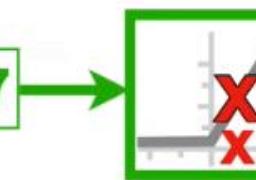
Filter
(Convolution)



Input to NN



sum $+ 0.97$



$\times -1.33$

$\times 1.33$

$+ -0.45$

$+ 1.45 =$

1

90

~~90~~



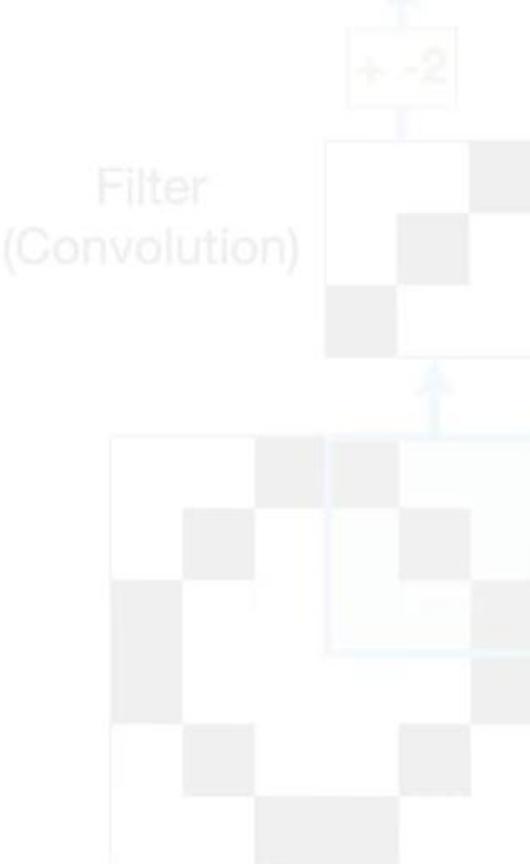
Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1

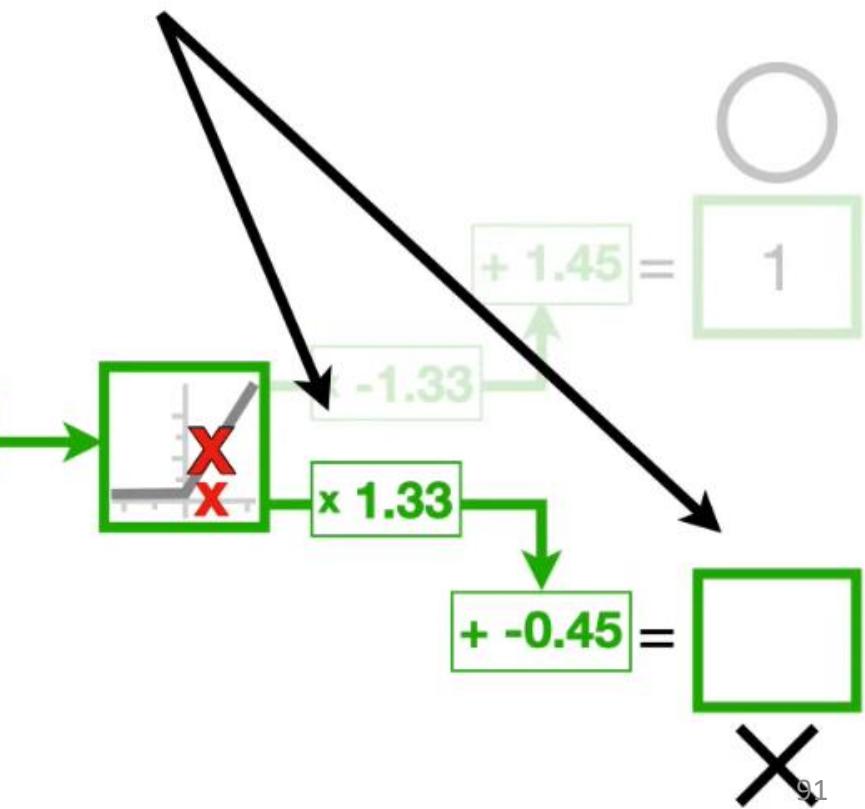
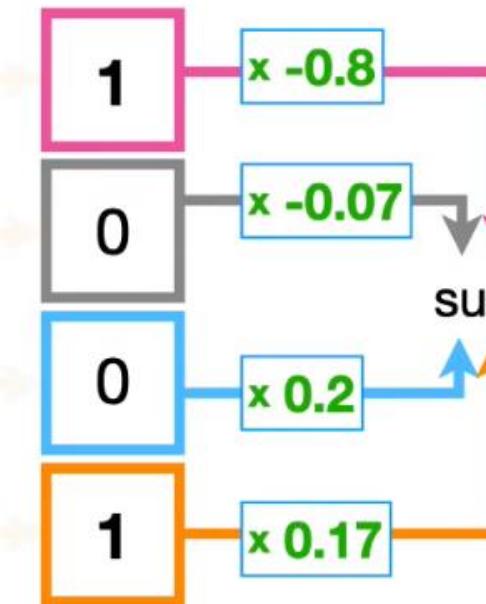
And the connection from the **Hidden Layer** to the output for the letter “X”...



Input to NN

Max Pooling

1	0
0	1





Feature Map

1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1

$$+ -2$$



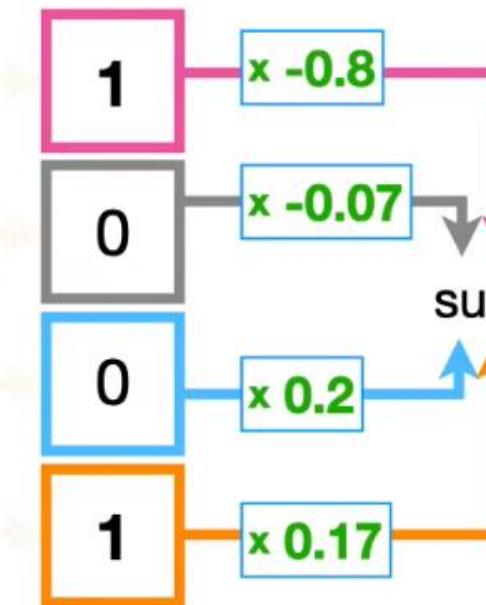
Filter
(Convolution)



1	0
0	1

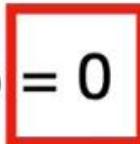
Max Pooling

Input to
NN



...gives us **0** for the letter "**X**".

$$(0.34 \times 1.33) + -0.45 = 0$$





Feature Map

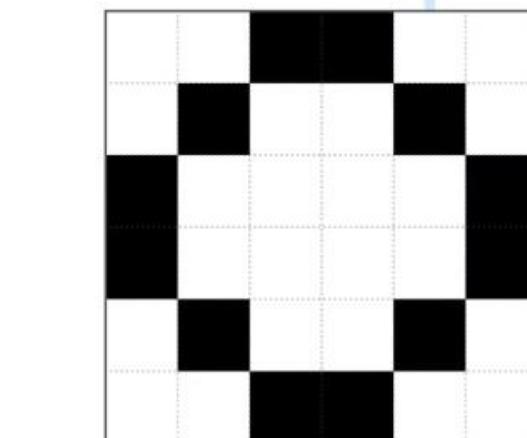
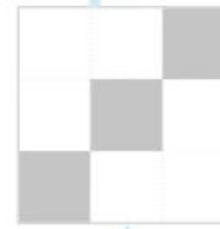
1	-1	-2	-1
-1	-2	-1	-2
-2	-1	-2	-1
-1	-2	-1	1



1	0	0	0
0	0	0	0
0	0	0	0
0	0	0	1

So, when the input is a picture of the letter “O”...

Filter
(Convolution)



Input to NN

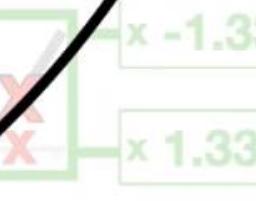
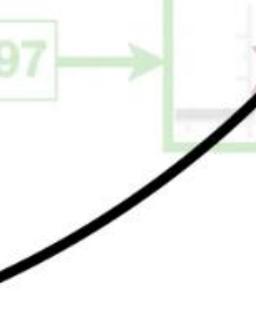
NN

1	0	
0	1	

1	x -0.8
0	x -0.07
0	x 0.2
1	x 0.17

sum

$$+ 0.97$$



$$+ 1.45 =$$

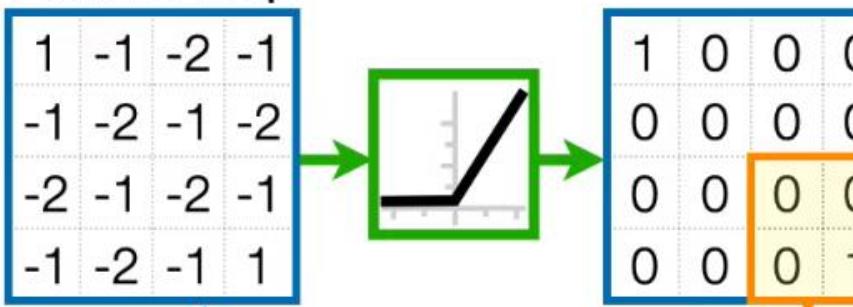
$$1$$

$$+ -0.45 =$$

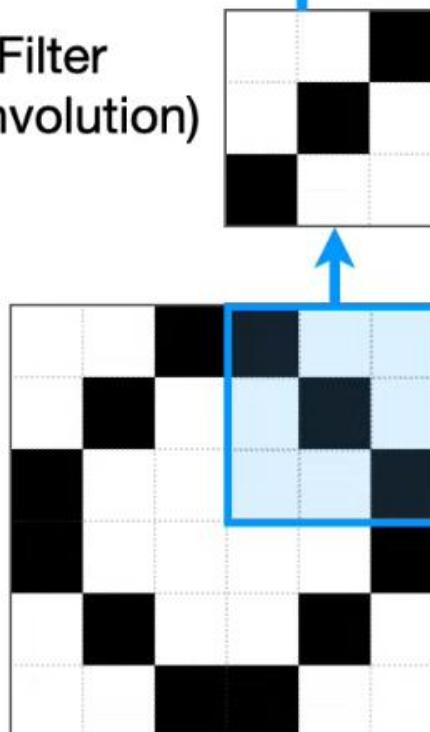
$$0$$



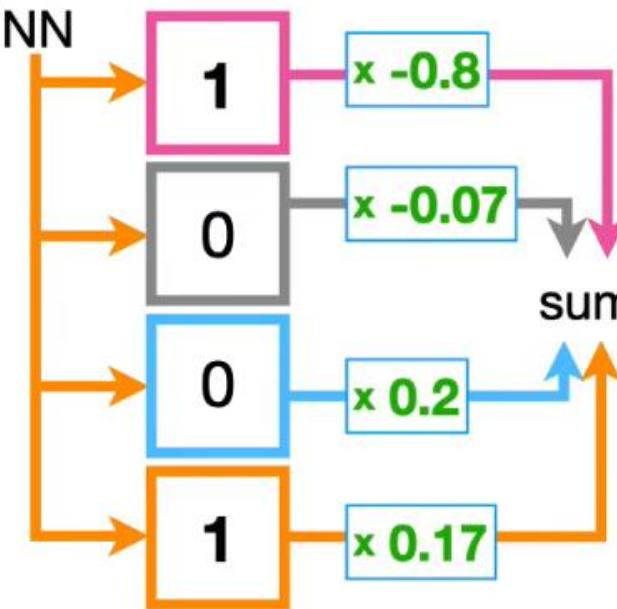
Feature Map



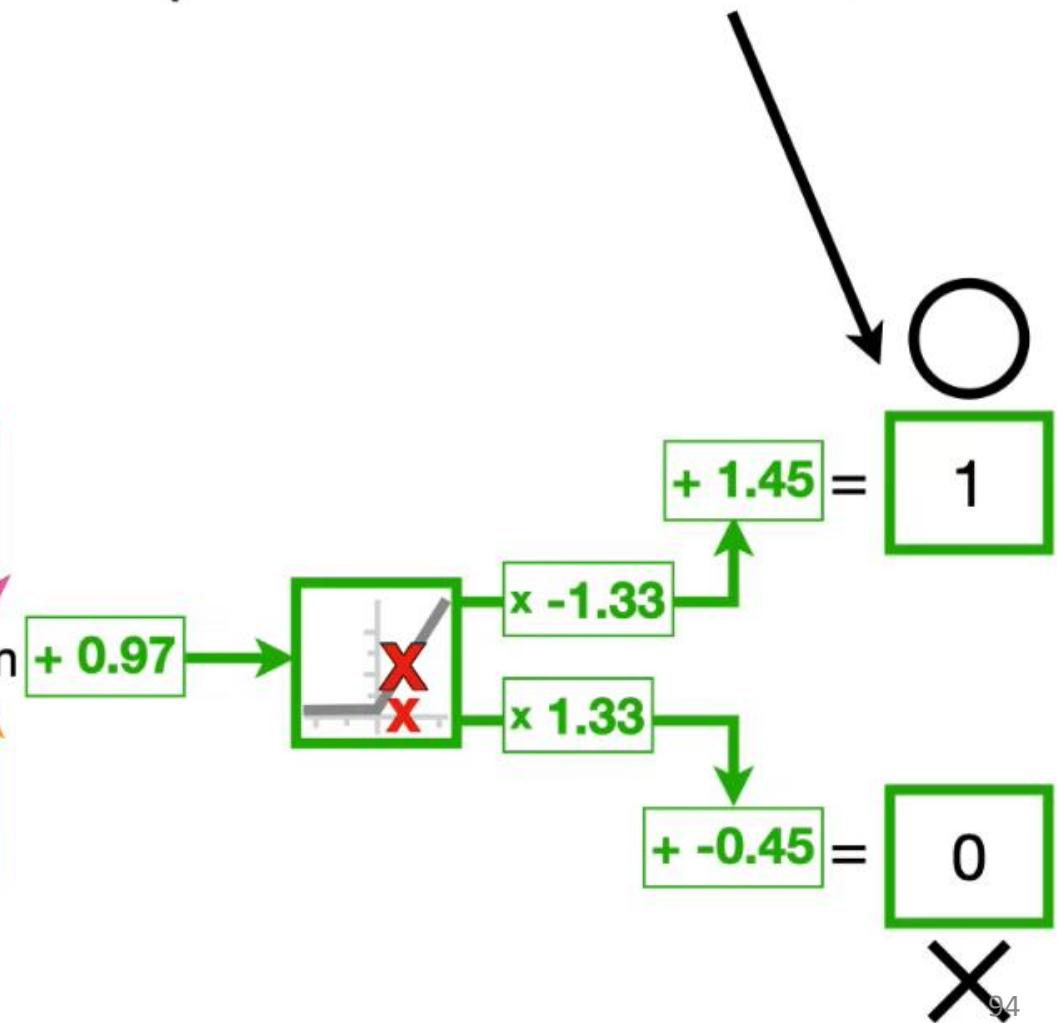
Filter
(Convolution)



Input to
NN



...this **Convolutional Neural Network** classifies it as a picture of the letter “O”.



Summary:

- Convolution Layer → Feature Map.
- Max-Pooling Layer → Max Value from group of Neighboring values.
- Fully Connected layer(NN) → Understand high level features extracted by Previous layers.

About Time Complexity...

Time complexity is not a crucial aspect of machine learning and neural networks and it's not always the primary focus of discussion for a few reasons

- **Empirical Emphasis**
- **Algorithmic Complexity**
- **Complexity Trade-offs**
- **Parallelization and Distributed Computing**

Time Complexity:

$$O(N^2 * K^2 * C_{in} * C_{out} + N^2 * C_{out} * P^2 + L * M^2)$$

- Input feature map size: $N \times N$
- Kernel size: $K \times K$
- Number of input channels: C_{in}
- Number of output channels (filters): C_{out}
 $O(N^2 * K^2 * C_{in} * C_{out})$
- Input feature map size: $N \times N$
- Number of input channels: C_{out} (from the previous layer)
- Pooling window size: $P \times P$
 $O(N^2 * C_{out} * P^2)$
- Assuming you have L fully connected layers with the following parameters:
- Number of neurons in each fully connected layer: M
 $O(L * M^2)$

THANK YOU !